

CSE 127 Computer Security

Stefan Savage, Spring 2020, Lecture 8

System Security II: Side Channels

Context

- Last class we talked about using **isolation** and **privilege levels**
 - Used to implement privilege separation, least privilege and complete mediation
 - Basic idea: **protect the secret or sensitive stuff so it can't be accessed across a trust boundary**
- Assumption: we know what the trust boundaries are and that **access** to something is easy to identify

Side Channels

- We often think of systems as black boxes:
 - As abstractions that consume input and produce output
 - We assume that all side effects are about output (e.g., values in memory or I/O)
- Sometimes, in addition to what is in the output, critical information can be revealed in **how** it is produced.
 - Yes, how... like how long, how fast, how loud, how hot... artifacts of the **implementation** not the **abstraction**
 - This can produce a **side channel**: a source of information beyond the output specified by the abstraction.

Today

- Overview of side channels in general
- Cache side channels
- Meltdown, Spectre & Rowhammer

Consumption Side Channels

- How long does this password check take?
 - Depends on where the first mismatch is.


```
char pwd[] = "z2n34uzbnqhw4i";  
  
//...  
  
int CheckPassword(char *buf)  
{  
    return strcmp(buf, pwd);  
}
```

Side Channels

- Consumption: how much of a resource is being utilized to perform the operation?
 - Examples: time, power, memory, network, etc.
- Emission: what out-of-band signal is generated in the course of performing the operation?
 - Examples: electro-magnetic radiation, sound, movement, error messages, etc.

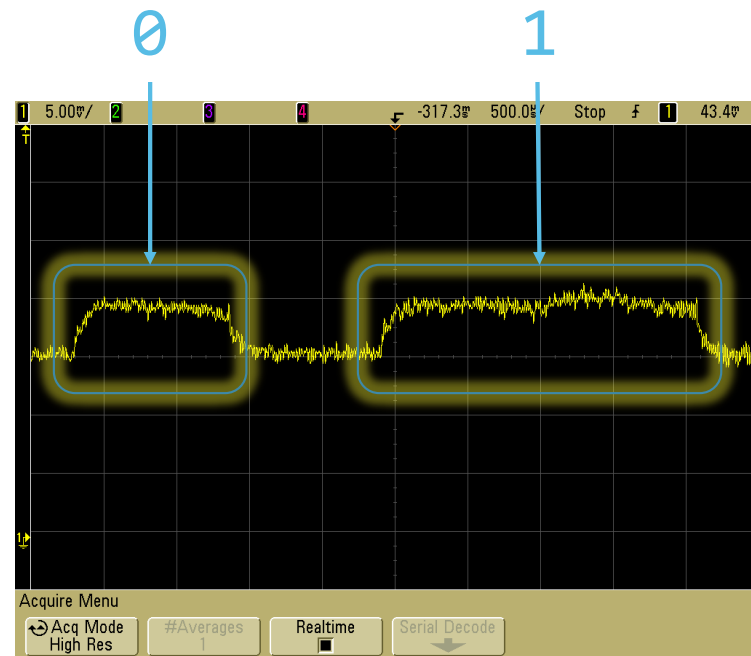
Side Channel Examples

- Tenex password verification
 - Alan Bell, 1974
 - Character-at-a-time comparison + virtual memory
 - Recover the full password in linear time



Side Channel Examples

- Secret cryptographic key value maintained in hardware
 - Can never be read, only used
 - (e.g., cable box, smart phone)
- Simple Power Analysis (SPA)
- Differential Power Analysis (DPA)
 - Paul Kocher, 1999
 - Using signal processing techniques on a very large number of samples, iteratively test hypothesis about secret key bit values.



Side Channel Examples

- Timing Analysis of Keystrokes and Timing Attacks on SSH
 - Dawn Song, David Wagner, Xuqing Tian, 2001
 - Recover characters typed over SSH by observing packet timing




Figure 1: The traffic signature associated with running SU in a SSH session. The numbers in the figure are the size (in bytes) of the corresponding packet payloads.

Side Channel Examples

- Side-Channel Leaks in Web Applications: a Reality Today, a Challenge Tomorrow
 - Shuo Chen, Rui Wang, XiaoFeng Wang, Kehuan Zhang, 2010
 - Recover characters typed into search boxes over HTTPS by observing sizes of responses

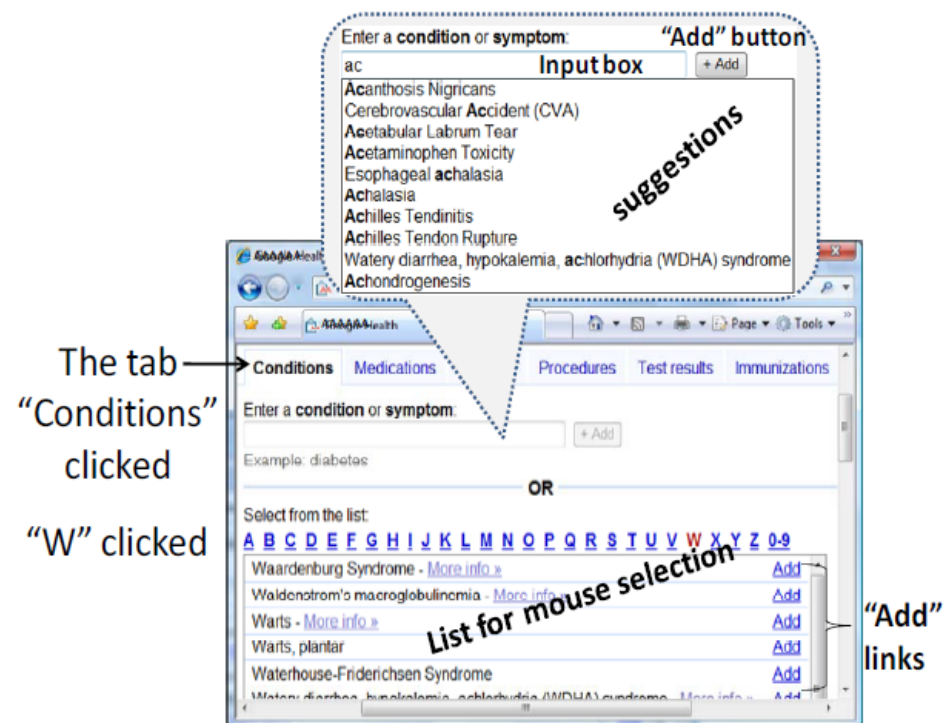


Figure 2: User interface for adding health records

Side Channel Examples

- Keyboard Acoustic Emanations
 - D. Asonov, R. Agrawal, 2004
 - Recover keys typed by their sound
- Keyboard Acoustic Emanations Revisited
 - Li Zhuang, Feng Zhou, J. D. Tygar, 2009


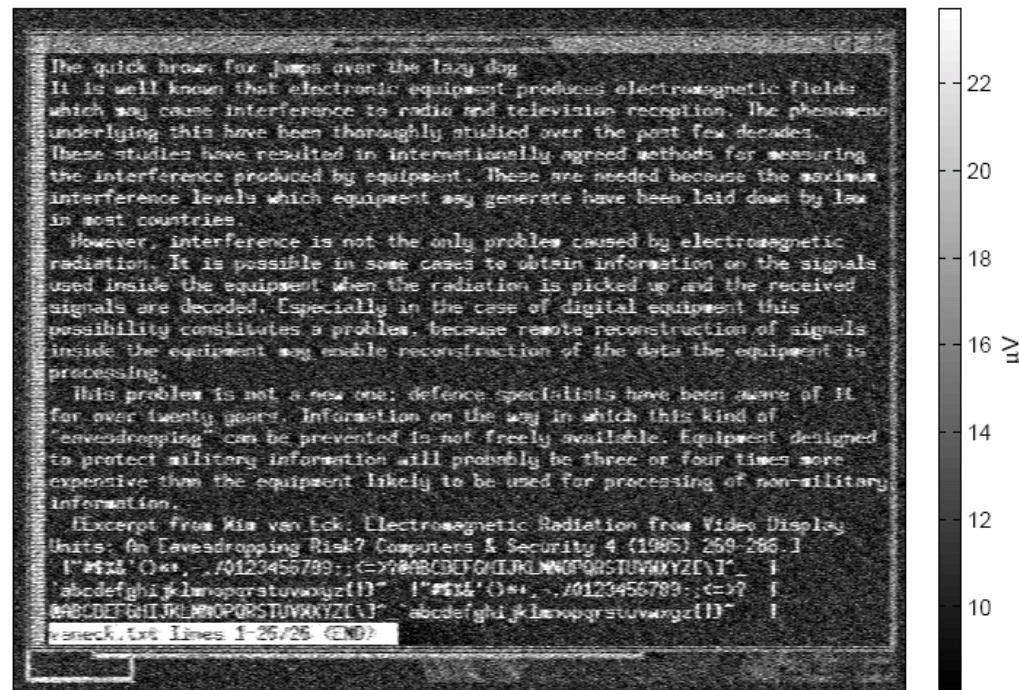


Fig. 7. Length of recording vs. recognition rate.

Remote reading of LCD screens via RF (Kuhn, 2004)

350 MHz, 50 MHz BW, 12 frames (160 ms) averaged



Target and antenna in a modern office building 10 m apart, with two other offices and three plasterboard walls (-2.7 dB each) in between. Single-shot recording of 8 megasamples with storage oscilloscope at 50 Msamples/s, then offline correlation and averaging of 12 frames.

Optical domain emanations (Kuhn, 2002)

- Light emitted by CRT is
 - Video signal combined with phosphor response
- Can use fast photosensor to separate signal from HF components of light
- Even if *reflected* off diffuse surface (i.e., a white wall) from across the street


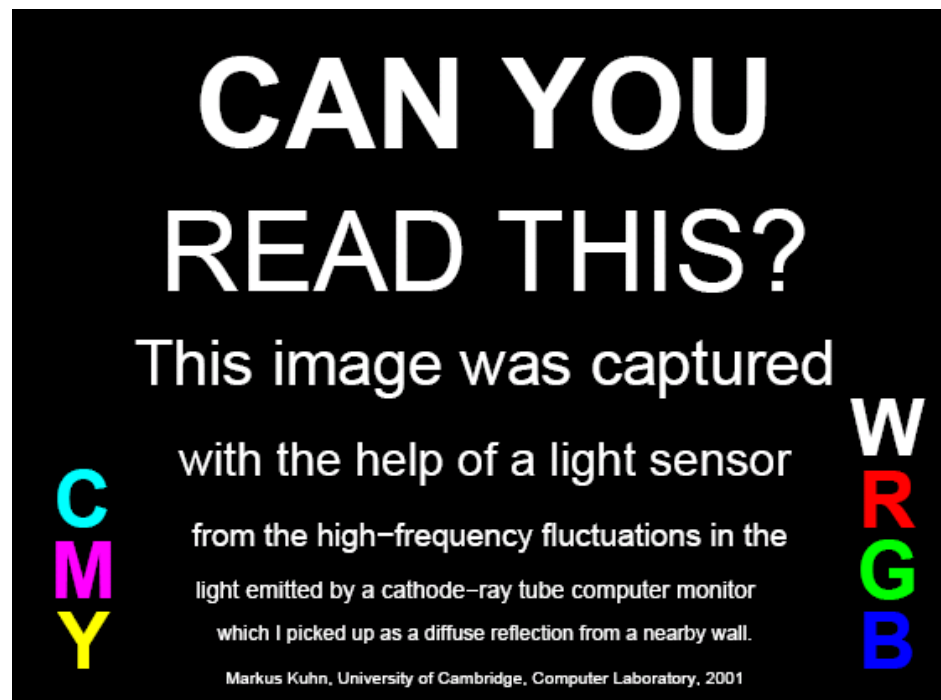


Figure 1. Photomultiplier tube module.

Source signal



Bounced off a wall



Side Channel examples

- Heat of the Moment, 2011
 - Meiklejohn et al




Figure 2: The Diebold plastic ATM keypad with rubber keys, model 19-019062-001M REV1.

Side Channels

- The Pentagon Press Parking Index

[Follow](#)

I hesitate to post the Pentagon Press Parking Index (PPPI) tonight because if it's flashing red tomorrow, I may not share. But for the next 12 hours, prospects of Assad strikes are low.



4:01 PM - 11 Apr 2018

<https://twitter.com/HansNichols/status/984204920722264064>

Active Side Channels

- Faults can create additional side channels or amplify existing ones
 - Erroneous bit flips during secret operations may make it easier to recover secret internal state
- Attackers can induce faults, leading to ***fault injection attacks***
 - Glitch power, voltage, clock
 - Vary temperature
 - Subject to light, EM radiation

Aside: covert channels

- Side channels are inadvertent artifacts of the implementation that can be analyzed to extract information across a trust boundary
- **Covert channels** are the same idea, but put on purpose
 - One party is trying to leak information in a way that it won't be obvious
 - By encoding that information into some side channel
 - E.g., variation in time, memory usage, etc...
 - Incredibly difficult to protect against

Mitigating Side Channels

- Eliminate dependency on secret data
- Make everything the same
 - Use the same of amount of resources every time
 - Hard (many optimizations in hardware, compilers, etc.)
 - Expensive (everything runs at worst-case performance)
- Hide
 - “Blinding” can be applied to input for some algorithms
 - Secret data is not operated on in “clear” form
 - Only possible for certain algorithms (eg RSA)
- Adding random noise?
 - Attacker just needs more measurements to extract signal

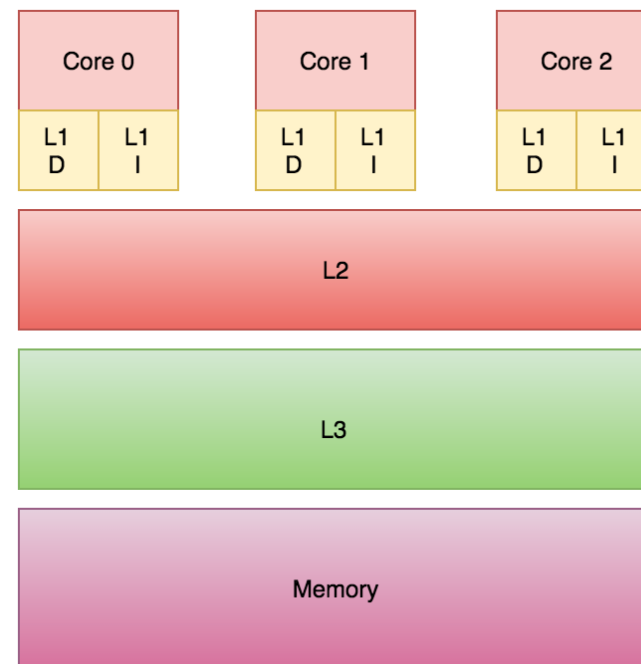
Cache Side Channels

Cache

- Main memory is dense (high capacity) ... but slow
- Processors try to “**cache**” recently used memory in faster, but smaller capacity, memory cells closer to the actual processing core



Cache Hierarchy

- Caches are such a great idea, let's have caches for caches!
- Level closest to the processing core is fastest but smallest capacity
- Successive levels have increasing capacity, but slower speeds



Cache Organization

- Cache line
 - Unit of cache granularity
 - Example: 64 bytes
- Set Associativity
 - Cache lines grouped into sets
 - Each memory address is mapped to a set of cache lines
 - Ex: $\text{address} \% \text{number_of_sets}$
 - Aside: which address? (typically)
 - L1: Virtually indexed
 - L2/L3: Physically indexed
 - All typically “tagged” by physical address
 - Can use either set to satisfy (reduces conflict misses)



Cache Side Channels

- Cache is a shared system resource
 - Not isolated by process, VM, or privilege level
 - “Just a performance optimization”
 - Has no impact on *contents* of reads
 - But does impact the *time* that a read takes

Cache Side Channels

- Threat model:
 - Attacker and victim are two different execution domains (processes or privilege levels) on the same physical system
 - Attacker is able to invoke (directly or indirectly) functionality exposed by the victim
 - Potentially, with attacker-supplied parameters
 - Example: using API, making a system call, running a setuid binary, etc.
 - We don't want attacker to be able to infer anything about the contents of victim memory
- Some algorithms, including cryptographic algorithms, have memory access patterns that are dependent on sensitive memory contents
 - Therefore memory access patterns may implicitly leak memory contents

Things attacker can do

- Prime
 - Place a known address in the cache (by reading it)
- Evict
 - Access memory until a given address is no longer cached (force capacity misses)
- Flush
 - Remove a given address from the cache (`cflush` instruction on x86)
- Measure
 - Very precisely (down the the cycle) how long it takes to do something (`rdtsc` on x86)

All cache side channel attacks of a similar form:
manipulate cache into known state, make victim run and then try to infer what has changed in cache as a result

Cache Side Channels

- Three basic techniques
 - **Evict and time**
 - Kick stuff out of the cache and see if victim slows down as a result
 - **Prime and probe**
 - Put stuff in the cache, run the victim and see if you slow down as a result
 - **Flush and reload**
 - Flush a particular line from the cache, run the victim and see if your accesses are still fast as a result

Evict & Time

- Run the victim code several times and time it
 - Baseline
- Evict (portions of) the cache
 - How?
 - Access many memory locations so that previous cache contents are replaced with this recently-accessed data (conflicts in cache cause eviction)
- Run the victim code again and retime it
- If it is slower than before, then cache lines evicted by the attacker must've been used by the victim
 - We now know something about the addresses accessed by victim code

Prime & Probe

- Prime the cache
 - Access many memory locations (covering all cache lines of interest) so previous cache contents are replaced with attacker addresses
 - Time access to each cache line to establish speed for “in cache” references
- Run victim code
- Attacker retimes access to its own memory locations
 - If any are slower then it means the corresponding cache line was used by the victim
 - We now know something about the addresses accessed by victim code

Flush & Reload

- Specifically for shared memory
 - E.g., shared libraries, fork() sharing, deduplication in VMs
- Time memory access to (potentially) shared regions
- Flush (specific lines from) the cache
- Invoke victim code
- Retime access to flushed addresses, if still fast was used by victim
 - Because we flushed it it should be slow, victim must have reloaded it
 - We now know something about the addresses accessed by victim code

Also

- The error on any individual measurement is high
 - So you repeat many times and use normal order statistics (e.g., median values) to filter outliers from pattern
- Can work on different levels of caches (L1 to L3)
 - Different time domains, just need to exclude effects of faster cache levels
- Can work on both instruction cache and data cache
 - These are commonly separate at L1
 - Measurements and primes are simply about running code sequences rather than accessing data values

Cache Side Channels

- Ok, if I suspend all disbelief I may be able to imagine how something like this is possible in theory, but how practical is it really?
- *"Our error-correcting and error-handling high-throughput covert channel can sustain transmission rates of more than 45 KBps on Amazon EC2."*
 - *Hello from the Other Side: SSH over Robust Cache Covert Channels in the Cloud*
 - by Clementine Maurice, Manuel Weber, Michael Schwarz, Lukas Giner, Daniel Gruss, Carlo Alberto Boano, Kay Romer, Stefan Mangard



However...

- So far, we've only been able to infer things about the *addresses* being used
 - In certain cases this is bad (e.g., address pattern can give info about what values might be if computation depends on them), can be used for covert channels, etc
 - But...
- We can't yet
 - Read **protected values**
 - Write **protected values**
- Coming up: DRAM disturbance attacks (writing protected memory), Speculation attacks (reading protected memory)
 - Rowhammer, Meltdown and Spectre

Rowhammer

Review: RAM



- Volatile memory: data retained only as long as power is on
 - As opposed to persistent memory, which retains data even without power (e.g. flash, magnetic disks)
- Static RAM (SRAM) vs Dynamic RAM (DRAM)
 - SRAM: retains bit value as long as power is on, without any additional refresh
 - Faster
 - Lower density
 - Higher cost
 - DRAM: requires periodic refresh to maintain stored value
 - Refresh about every 64 ms
 - Higher density (higher capacity)
 - Lower cost



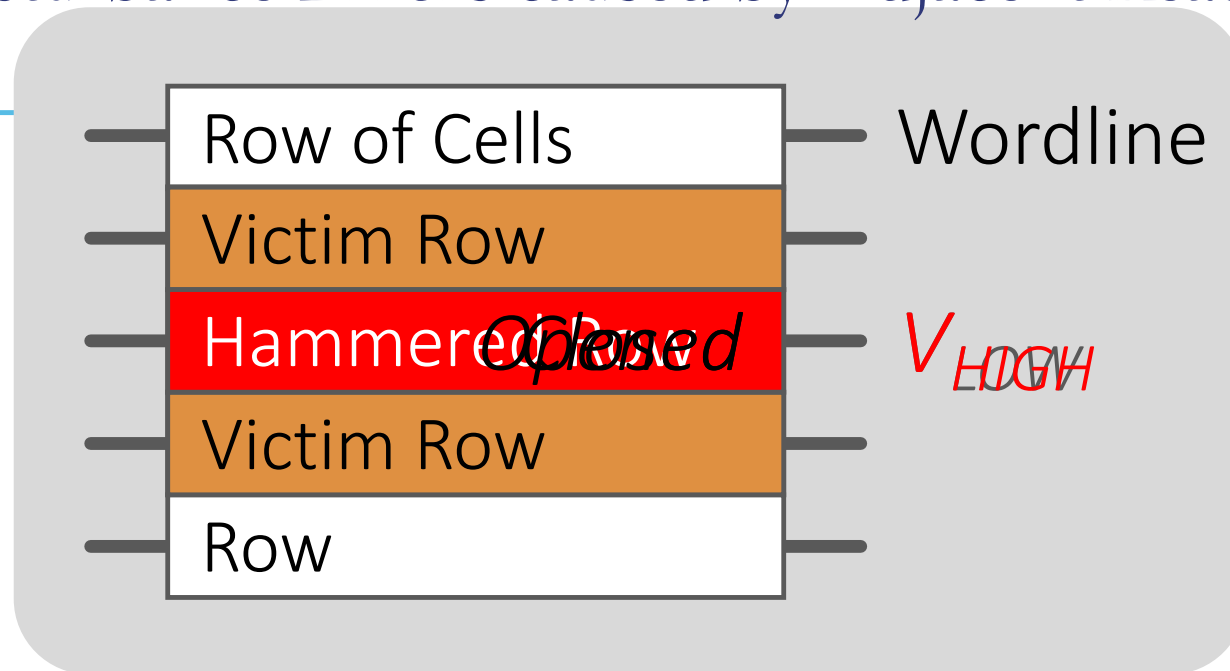
https://commons.wikimedia.org/wiki/File:SRAM_Cell_Inverter_Loop.png
<https://allthingsvlsi.wordpress.com/tag/dram-cell/>

Review: DRAM

- Cells are grouped into rows
 - ~1Kb per row
 - All cells in a row are refreshed together
- Refresh: read the row and write it back
- All access to individual cells happens via the “row buffer”
- As DRAM gets smaller (<35nm) more issues with *reliability* (retention, writes, and... reads)



Disturbance Errors caused by Adjacent Reads



Repeatedly opening and closing a row enough times within a refresh interval induces **disturbance errors** in adjacent rows in many of the real DRAM chips on the market

Rowhammer

- Attack scenario: attacker code is executing on the same machine as victim, but with less privileges
 - Example: userland attacking kernel, javascript attacking browser, etc.
 - Rowhammer attack lets attacker modify memory they can't access
- Exploit sketch
 - Characterize DDR flip locations
 - Identify protected target data to flip
 - Examples: page tables, su binaries, etc.
 - Maneuver protected data over flipping location
 - Example: allocate all other memory
 - Hammer own memory locations to alter protected data
- Amazing fact: **random** bit flips can be used to semi-deterministically take over computer

Rowhammer solutions

- ECC memory
 - Compute error correcting code on write, check on read
 - Significant mitigation, but recent attacks can bypass in some cases
 - Adds cost (common in servers, but not in consumer devices)
- Memory controller limitations on “hammering” or additional adjacent line refresh
 - Can have performance impact
- Process changes to mitigate inter-row effects
 - Conflicts with density demands




Meltdown and Spectre

Review: Virtual Memory

- Kernel virtual memory is mapped into every process
 - For efficiency
- Page table access control ensures that kernel pages are only accessible when the processor is in privileged mode

	Non-Secure				Secure	
EL0	App X	App Y	App X'	App Y'	App X''	App Y''
EL1	Guest OS A		Guest OS B		Secure OS	
EL2	Hypervisor					
EL3			Secure Monitor			



Quick review: ISA and μ Architecture

Instruction Set Architecture (ISA)



- Defined interface between hardware and software
 - Registers, instructions, etc.

μ Architecture

- Implementation of the ISA
- “Behind the curtain” details
 - E.g. cache specifics
- Key issue:: μ Architectural details can sometimes become “architecturally visible”

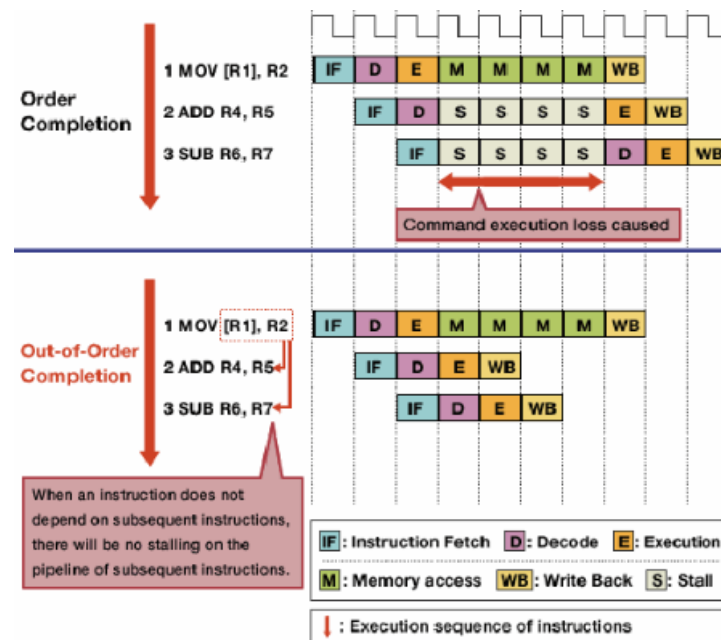
Review: Instruction Pipelining

- Processors break up instructions into smaller parts so that these parts could be processed in parallel.
- μ Architectural optimization
 - Architecturally, instructions appear to be executed one at a time, in order
 - Dependencies are resolved behind the scenes



Review: Out-of-order Execution

- Some instructions can be safely executed in a different order than they appear
- Avoid unnecessary pipeline stalls
- μ Architectural optimization
 - Architecturally, it appears that instructions are executed in order



Review: Speculative Execution

- Sometimes control flow depends on output of an earlier instruction.
 - E.g. conditional branch, function pointer
- Rather than wait to know for sure which way to go, the processor may “speculate” about the direction/target of a branch
 - Guess based on the past
 - If the guess is correct, performance is improved
 - If the guess is wrong, speculated computation is discarded and everything is re-computed using the correct value.
- μ Architectural optimization
 - At the architecture level, only correct, in-order execution is visible

Excerpted from

Embedded Security

Chapter 3: Meltdown and Spectre

Michael Schwarz, Moritz Lipp

20.03.2018

www.iaik.tugraz.at

Switch to Schwarz and Lipp slides

What to do?

- Meltdown

- Key issue is that privilege checks are not being checked before cache access allowing user inference of kernel memory contents
- Solution: stop mapping kernel memory into process address space
Makes system calls slower
(a range of optimizations, including AS identifiers in TLBs make it better)

- Spectre

- Not so easy... general issue about speculation and branch prediction
- Could disable speculation on branches... but huge performance impact
 - Also can only be done by chip manufacturer (microcode patch)
- Selectively insert instructions to stop speculation at sensitive branches
 - LFENCE, (x86), CSDB (ARM)
 - Various special hacks (e.g., Intel adds special microcode sequences to manipulate branch prediction state)
- Ultimately, change cache architecture to not reflect speculated actions... but need entirely new hardware

These are the beginning on a wave

- Computer architects spent the last 20 years optimizing for the common case
 - Assumption is that if optimization doesn't change output then its all good
- All of these optimizations are now being examined and being used to create new side channels
 - L1TF, MDS, TAA, iTLB multihit
 - And more...
- Hardware vendors and computer architects are retooling to figure out how to still offer optimization without big security holes

Summary

- Side channels are a way to violate trust boundary even if there are no errors in interface or control flow vulnerabilities
 - Information contained in details of the implementation
 - Resource consumption or emanations
- Cache side channels
 - Shared caches provide a timing channel to infer what addresses are being used by victim code
- Rowhammer
 - Active channel for manipulating DRAM via reads (based on imperfect memory reliability)
- Speculation side channels
 - Speculative instructions can impact cache state, allowing measurements of protected state

Additional Resources

- Google Project Zero
 - <https://googleprojectzero.blogspot.com/>
 - Rowhammer:
 - <https://googleprojectzero.blogspot.com/2015/03/exploiting-dram-rowhammer-bug-to-gain.html>
 - Spectre/Meltdown:
 - <https://googleprojectzero.blogspot.com/2018/01/reading-privileged-memory-with-side.html>
- Anders Fogh blogs
 - <https://cyber.wtf/author/andersfogh1974/>
 - <https://dreamsofastone.blogspot.com/>
- Paul Kocher on Spectre
 - https://www.youtube.com/watch?v=hqlavX_SCWc

Additional Resources

- TU Graz team
 - Moritz Lipp: <https://mlq.me/>
 - Daniel Gruss: <https://gruss.cc/>
 - Michael Shwarz: <https://misc0110.net/web/>
- VUSec
 - <https://www.vusec.net/>