

CSE 127 Computer Security

Stefan Savage, Spring 2020, Lecture 14


Network Security II: DNS, Denial-of-service and perimeter defense

Today

- The Domain Name System
 - Another place where names at different layers can bound
 - The target of quite a few attacks
- Denial of service
 - Network attacks on availability
- Network perimeter defenses
 - Firewalls, NATs, NIDS/NIPS

Remember this? How did this work?

- Where is www.yahoo.com?




Ignore for now:
How did you know the address of the Local DNS server?
How did you send a message to it?
How did the Local DNS server know the answer?

Domain Name System

- We humans do not tend to remember 32bit IP numbers...
- Solution: domain names
 - Human readable identifiers (e.g., www.cs.ucsd.edu, google.com, etc)
- Problem: how to map DNS names to IP addresses?
 - In the old days we had a big file – *literally* (download from sri-nic.arpa)
 - Today we use a distributed name servers called the Domain Name System (DNS)

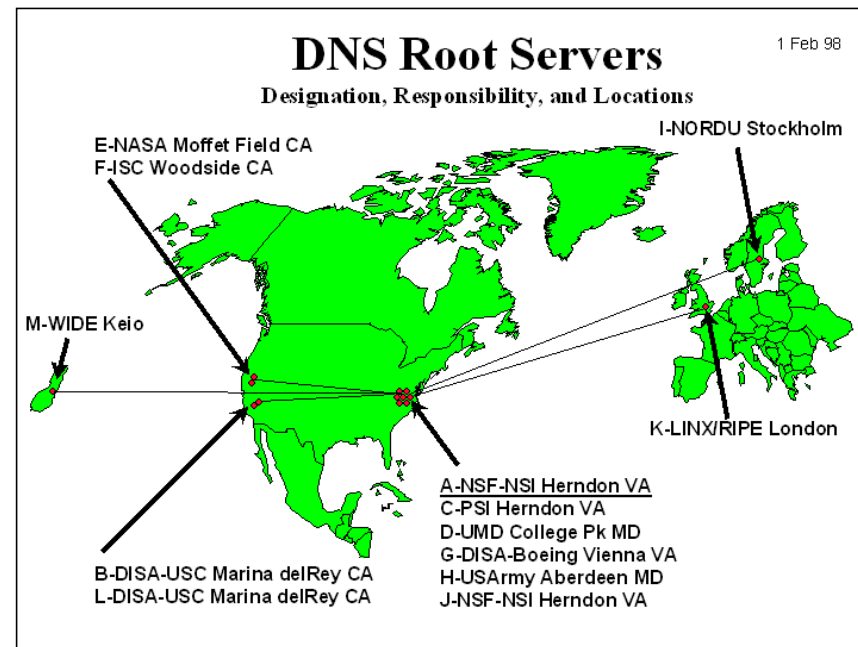
Domain Name System (DNS)

- Hierarchical Name Space

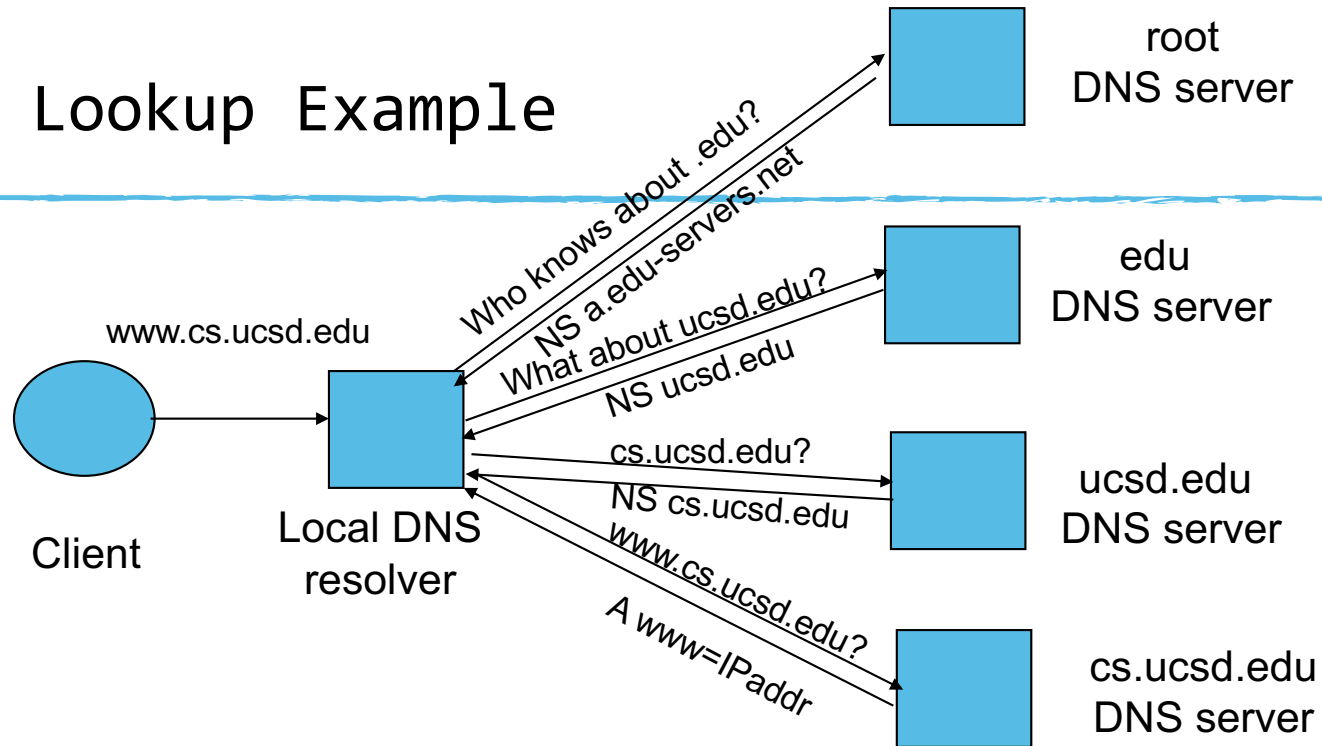


DNS Root Name Servers

- Hierarchical service
 - 13 root name servers for top-level domains
 - **Hardcoded into all systems**
 - Choose one at random
 - Authoritative name servers for subdomains
 - Local name resolvers (also called recursive resolvers) contact authoritative servers when they do not know a name



DNS Lookup Example



DNS record types (partial list):

- NS: name server (points to other server)
- A: address record (contains IP address)
- MX: address in charge of handling email
- TXT: generic text (e.g. used to distribute site public keys (DKIM))

Caching

- DNS **responses are cached**
 - Quick response for repeated translations
 - Useful for finding servers as well as addresses
 - NS records for domains
- DNS **negative queries** are cached
 - Save time for nonexistent sites, e.g. misspelling
- Cached data periodically times out
 - Lifetime (TTL) of data controlled by owner of data
 - TTL passed with every record, delete cached entry after TTL expires

DNS cache poisoning


- Basic idea:
 - If I can convince a DNS resolver to cache a bad mapping (e.g., www.cs.ucsd.edu points to 127.0.0.1) then everyone who uses it for www.cs.ucsd.edu will get that incorrect resolution
 - Can then be used for fraud, man-in-the-middle attacks, etc
- Used in lots of attacks
 - January 2005, the domain name for a large New York ISP, Panix, was hijacked to a site in Australia.
 - In November 2004, Google and Amazon users were sent to Med Network Inc., an online pharmacy
 - In March 2003, a group dubbed the "Freedom Cyber Force Militia" hijacked visitors to the Al-Jazeera Web site and presented them with the message "God Bless Our Troops"
 - 2000 campaign: Hilary2000.org -> hilaryno.com

But how to do it?

- Man-in-the-middle attacker: easy
 - Observe DNS requests from server
 - Send false response to server and block true response
- Passive attacked: easy also
 - Observe DNS requests from server
 - Send false response to server before true response
- What about off-path attacker
 - If you know when user is likely to make request, you can flood responses to their server blindly
 - But there's a problem – matching request and response

DNS Packet

- Query ID:
 - 16 bit random value
 - Links response to query




(from Steve Friedl)


Resolver to NS request

Context

- SBC Global customer looks up unixwiz.net
- Goes to their DNS server dnsr1.sbcglobal.net
- This is the request sent from that server to c.gtld-servers.net which can give answers for domains ending in .net




Response ignored if unrecognized QueryID



DNS additional section

- Answers to questions you didn't ask...
 - There is a good reason for it... if I tell you that the name server for foo.com is ns1.foo.com... how do you find its IP address?
- But this is a problem... what if I run a DNS server for foo.com and when I get asked for the IP address for bar.foo.com, I put some additional stuff in the "additional section" like:
 - You can find the IP address for paypal.com at 41.2.6.2 (address I control)
 - And you can also find the IP address for amazon.com and chase.com there too...

Early Attack Strategy

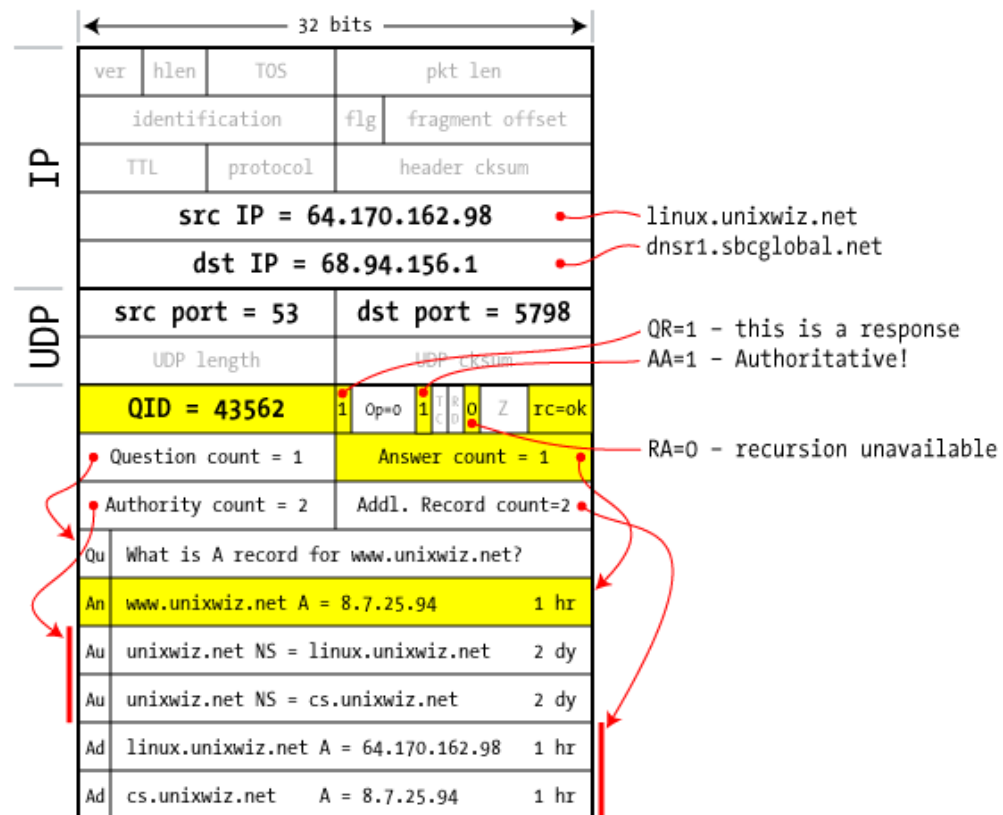


(so

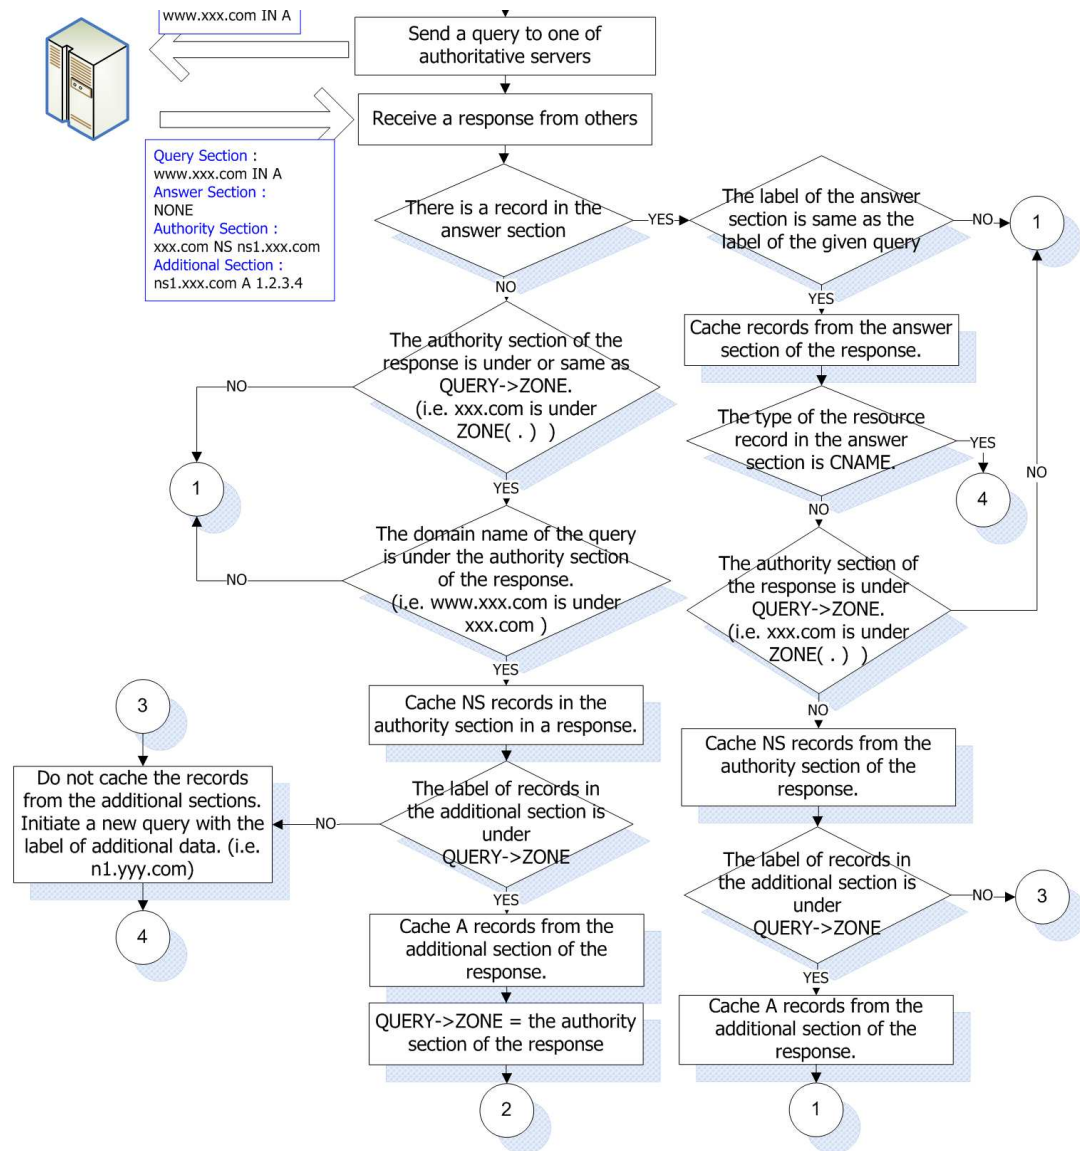
Authoritative response to resolver

bailiwick checking:
response is cached **only**
if it is within the same
domain of query
(i.e. a.com cannot
set NS for b.com)

final answer \rightarrow



Bailiwick Checking Rule from BIND



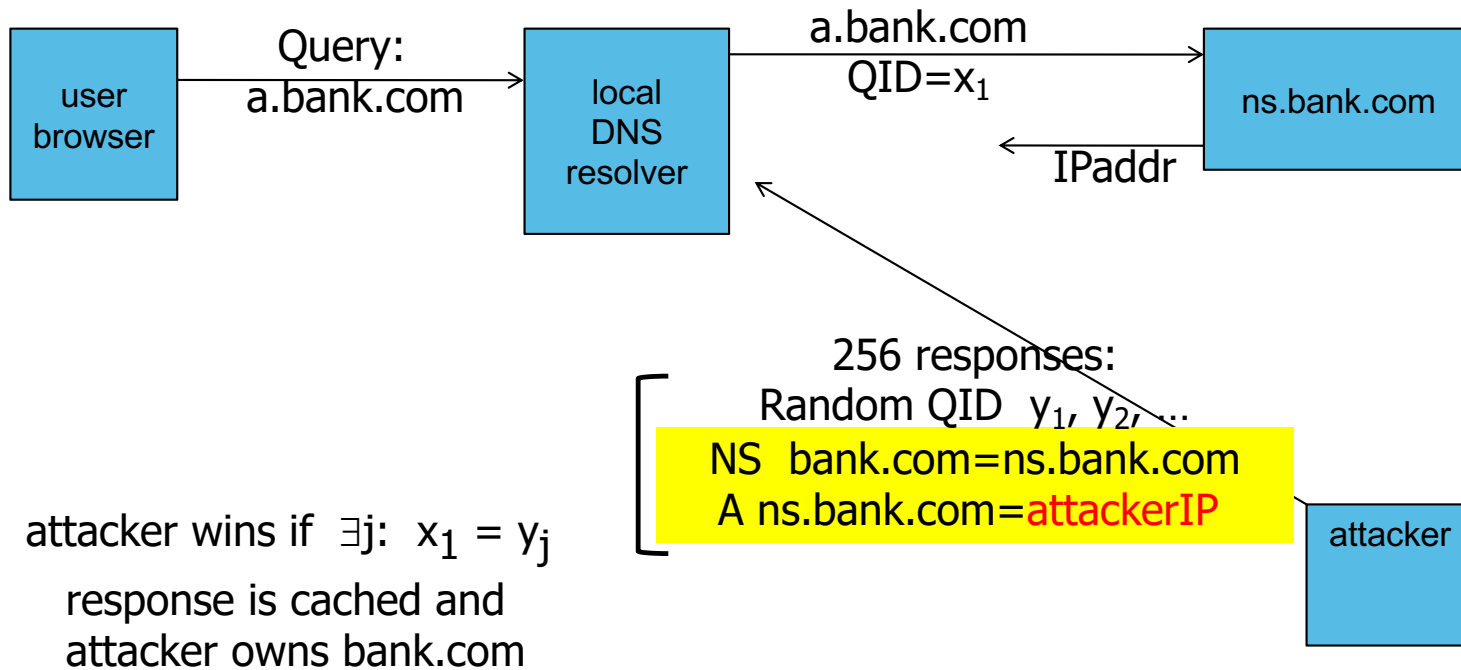
source: Son and Shmatikov, "The DNS Cache Poisoning" SECURECO

But we forgot something

- A decade goes by and Dan Kaminsky realizes that the bailiwick checking rule doesn't really protect us
- Unnoticed hole that allows arbitrary DNS poisoning at a distance
- To fix bug, *unprecedented* coordinated global operation to do secret mass migration of DNS infrastructure (2008)

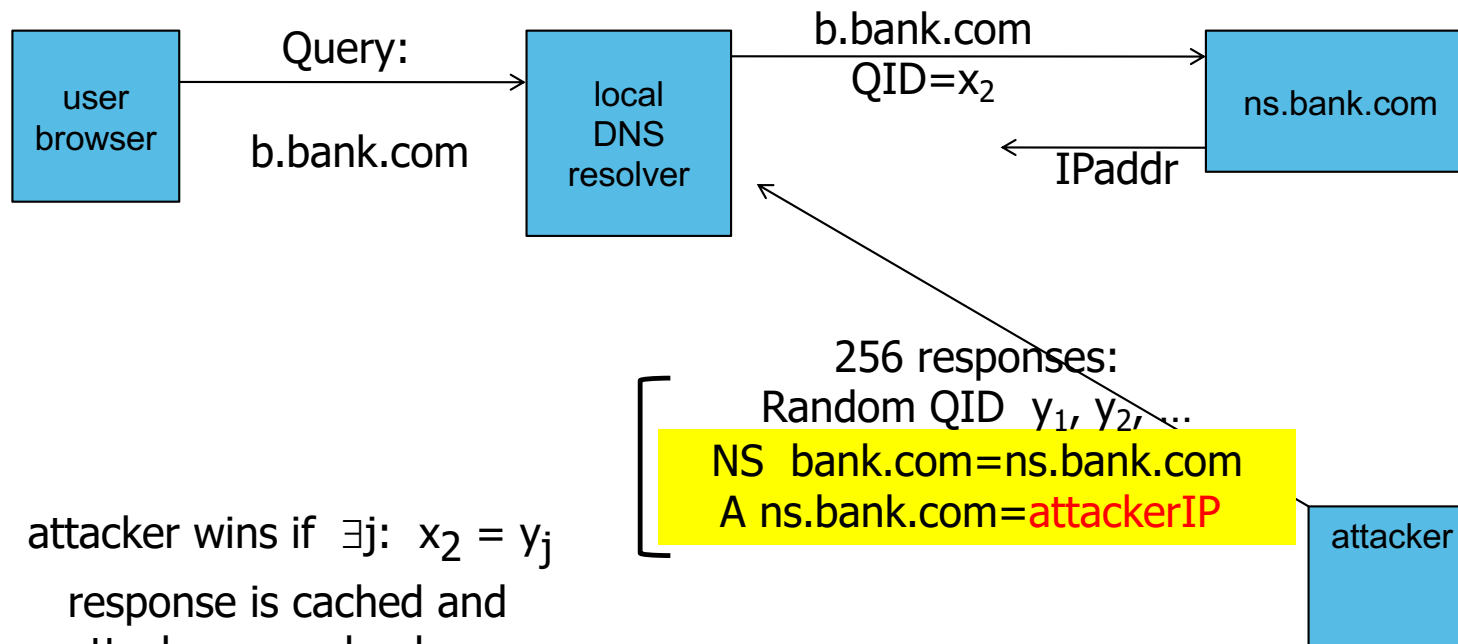
DNS cache poisoning (a la Kaminsky'08)

- Victim machine visits attacker's web site, downloads Javascript



If at first you don't succeed ...

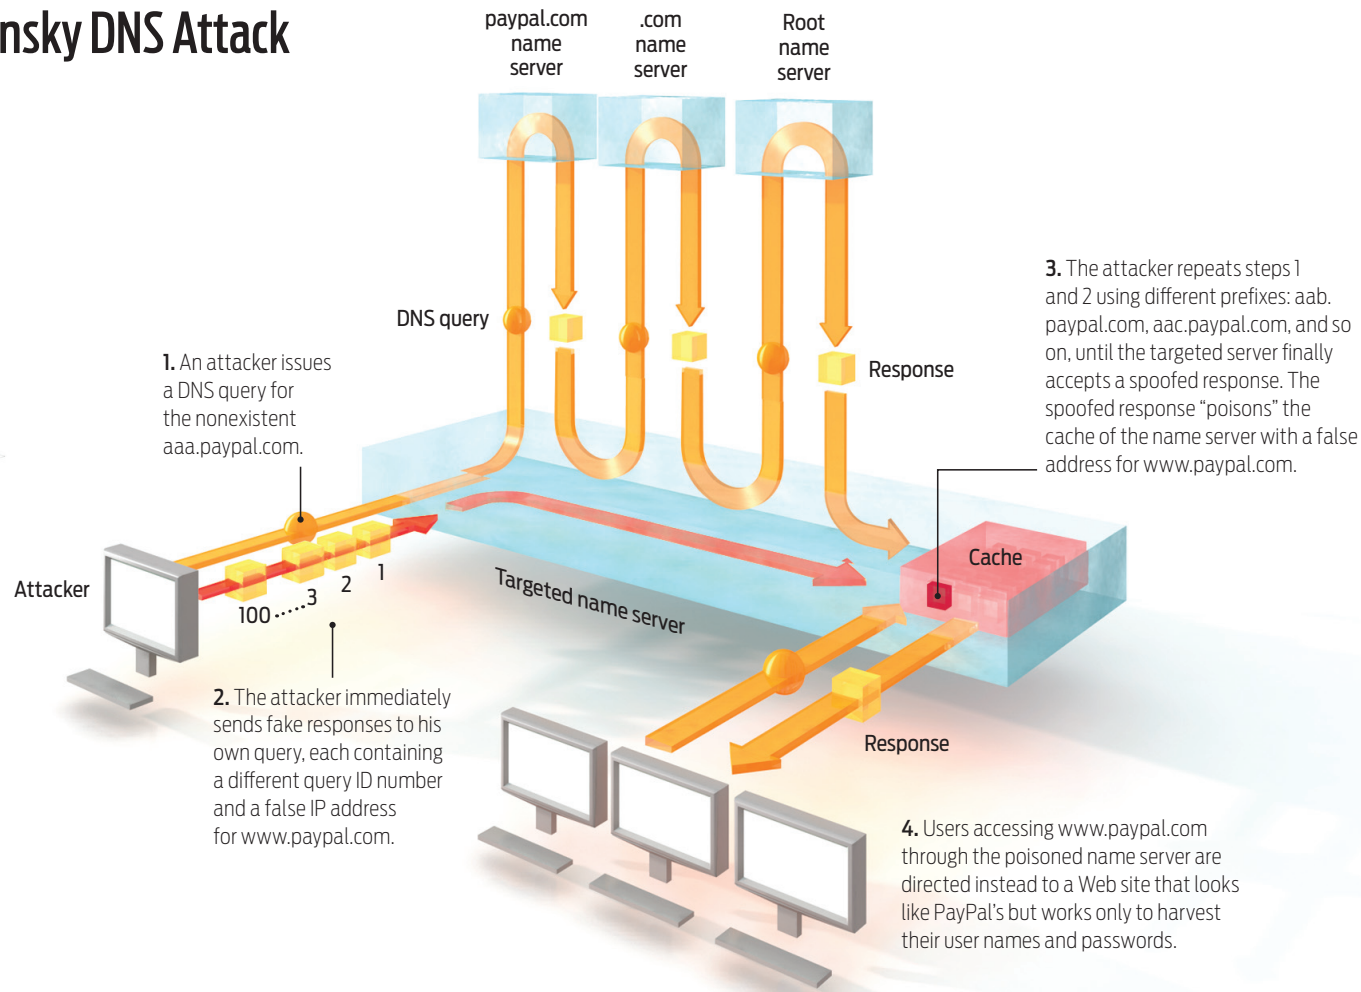
- Victim machine visits attacker's web site, downloads Javascript



attacker wins if $\exists j: x_2 = y_j$
response is cached and
attacker owns bank.com

success after ≈ 256 tries (few minutes)

Kaminsky DNS Attack



BRYAN CHRISTIE DESIGN

Defenses

- **Increase Query ID size.** How? Some approaches in use
 - Randomize src port and match it, additional 11 bits
 - Now attack takes many hours
 - Ox20 encoding – randomly vary capitalization (DNS is case insensitive) check that you get same capitalization back
- Try to detect poisoning
 - Ignore responses not directly necessary to query
- Authenticated requests/responses
 - Provided by DNSsec (digital signatures on DNS records) ... but few domains use DNSsec

DNS Summary

- Current DNS system does not provide strong evidence binding request to response
- Response can provide more data than was asked for
- Together allows attacker to “poison” DNS and divert traffic to their sites
- This is also why its so important for HTTPS to check certificate signatures (i.e., because just because you ask for www.amazon.com doesn't mean you'll get it)

Denial-of-service


- Attack against *availability*, not confidentiality, integrity, authenticity, etc
- Two kinds of attacks:
 - **Logic vulnerabilities:** exploit bugs to cause crash
 - e.g. Ping-of-Death, Land
 - Fix via filtering and patching
 - **Resource consumption:** overwhelm with spurious requests
 - e.g. SYN flood, bandwidth overflow
 - Much tougher to fix...
 - Same idea, more quantity: Distributed denial-of-service attacks (DDOS)
 - **Lots of hosts** attack a victim at once

Resource consumption of Service

- Server CPU/Memory resources
 - Consumes connection state (e.g. SYN flood)
 - Time to evaluate messages (interrupt livelock)
 - Some messages take “slow path” (e.g. invalid ACK)
 - Can cause new connections to be dropped and existing connections to time-out
 - Make DB process lots of queries
 - Attack cache – lots of random queries
- Network resources
 - Some routers are packet-per-second limited, FIFO queuing (send small pkts)
 - If attack is greater than forwarding capacity, good data will be dropped (send big pkts)

This is a widespread problem

- Most attacks small and focused on individuals or small sites
 - Particularly gamers; so-called booter services
- But some attacks are huge (e.g., 2018 Attack on GitHub)



What to do?

- Defenses against address spoofing
 - For attacks from randomly spoofed addresses
- Filtering based on attack features or IP address
- Buy more resources

Address spoofing

- Filter packets with incorrect source addresses
 - **Network egress**: filter outbound packets on a link whose source addresses are not reached using the link as the next hop (i.e., this couldn't be your source address)
 - But requires network routers to do "good deeds" for others...
 - **Network ingress**: filter inbound packets whose source address are not in the routing table at all
- SYN Cookies
 - Issue: allocating per TCP session state is **expensive** (that's why the SYN flood attack works)
 - Delay allocation of state until remote host commits to three-way handshake
 - Send back SYN/ACK packet **without allocating state** on server; server's initial sequence number (ISN) encodes a secret "cookie" that is function of some combination of (src,dst,srcport,dstport and time).
 - Allocate state when client sends ACK to server's SYN/ACK (using cookie to validate)

Address spoofing(2)

- TTL-based IP filtering
 - From a given host the TTL is decremented by a certain number of hops (based on network topology)
 - Std IP implementations set the packet TTL value to a small set of values (32, 64, 128, 255) [can normalize because Internet diameter is mostly < 32]
 - Thus, keep track of TTLs for each source network and if attack starts, filter packets whose TTLs are inconsistent
 - Example: suppose packets arriving from 128.2.x.x typically take 4 hops to get to you
So you'd expect TTLs of 28, 60, 124, or 251... maybe +/- 1
Attacker doesn't know what TTL to put in DoS packets sent to you
Discard packets from 128.2.x.x that have any other ttl values

Packet filtering

- Idea, if there is a common feature to the packet (i.e. “Die, you loser” in the payload, static source port #, odd TCP flags, etc.) then look for those packets and drop them
- If no feature exists then try to find way to add a “good” feature to legitimate flows that complete a 3-way handshake
- Instead of dropping packets, can simply rate-limit packets that are suspicious
- Third-party services will offer these capabilities on your behalf in the network
 - Like dialysis for Internet traffic – route network traffic to devices designed to filter out bad packets using one of the above techniques

Buy more resources

- Large content distribution networks (e.g. Akamai, CloudFlare) can handle very large attacks
- Each attacker gets diverted (i.e., via DNS) to local Akamai server instead of target
 - Total bandwidth Akamai can handle is the product of the bandwidth to all Akamai servers
 - Akamai has weathered attacks well in excess of 1TB/s
- Issue: who pays for that? \$\$\$

Special case: Reflection attacks

- Spoof source address to be that of victim
- Common example
 - Send name server request to 1000s of DNS servers *on behalf* of victim
 - All name servers send responses to victim
- Advantages (for attacker)
 - Amplification: some protocols response size \gg request size (NTP, DNS, SSDP)
 - Anonymity: attack doesn't come from attacker's machines
- Solution: try not to have "open" Internet services that allow this

DoS Summary

- In general, some of the toughest problems to solve
 - Network service model allows unsolicited requests
 - Bad guys can leverage large # of resources
 - Hard to attribute network actions
 - Few systems can account for effort spent per request or isolate impact of some requests from others
- DDoS-based extortion and retribution (e.g., against security companies) is not uncommon
Also widely used for political disruption

Network Perimeter Defense

- Idea: network defenses on “outside” of organization (e.g., between org and Internet)
 - Assumptions?
- Typical elements
 - Firewalls
 - Network Address Translation
 - Network Intrusion Detection/Prevention Systems(NIDS/NIPS)
(and other kinds of network content analysis)

Firewalls


- Access control policy
 - You'd like to limit who can talk to what service
 - Deny access to bad people; allow it to good
- Firewalls: One place to try to enforce this policy is **in the network**
 - Inbound: requests from outside the network to connect to services inside
 - Outbound: requests from inside the network to connect to services outside
- Conceptually simple idea motivates much of this
 - The outside is bad and scary, while people inside are good and friendly
 - Allow internal users to communicate with any outside service
 - Limit outside connections to a small set of services designed to be externally visible
- Questions:
 - What information do you use to filter?
 - Where do you do the filtering?

Kinds of Firewalls

- Personal firewalls
 - Run at the end hosts
 - e.g. Norton, Windows, etc.
 - Benefit: has more application/user specific information
- Network firewalls
 - Intercept and evaluate communications from many hosts
 - Deployed “in-line” in network infrastructure (i.e., mediates all communications)
 - Different levels of abstraction: packet filtering vs application proxies

Network Firewall

common deployment strategy



- Filters protect against “bad” communications.
- Protect services offered internally from outside access.
- Provide outside services to hosts located inside.

Packet Filtering Firewalls

- Packet filtering firewalls can take advantage of the following information from network and transport layer headers:
 - Source IP
 - Destination IP
 - Source Port
 - Destination Port
 - Flags (e.g. ACK)
- Some firewalls keep state about open TCP connections
 - Allows conditional filtering rules of the form “if internal machine has established the TCP connection, permit inbound reply packets”


Ports

- Ports are used to distinguish applications and services on a machine.
- Low numbered ports are often reserved for server listening.
- High numbered ports are often assigned for client requests.
- Imperfect language...
- Result: everything gets “tunneled” through ports not blocked by the firewall (e.g., 80, 443)
- Port 20 (TCP): FTP data
- Port 21 (TCP): FTP control
- Port 22 (TCP): ssh
- Port 25 (TCP): SMTP (Mail)
- Port 80 (TCP): HTTP
- Port 123 (UDP): NTP
- Port 143 (TCP): IMAP
- Port 443 (TCP): HTTPS
- Port 2049 (UDP): NFS
- Ports 6000 to 6xxx (TCP): X11

Principles for Firewall Configuration

- Least Privilege:
 - Turn off everything that is unnecessary
(e.g. Web Servers should disable port 25 [SMTP])
- Failsafe Defaults:
 - By default should reject
 - (Note that this can cause usability problems...)
- Egress Filtering:
 - Filter outgoing packets too!
 - You know the valid IP addresses for machines internal to the network, so drop those that aren't valid.
 - As per earlier, this can help prevent DoS attacks in the Internet.

Proxy-based Firewalls



- Proxy acts like *both* a client and a server.
 - Note, must terminate connection. Complications for HTTPS/TLS... must take out certificates
- More semantics available: able to filter using application-level info
 - For example, block based on URL, or on particular javascript strings
- Proxies can provide other services too
 - Caching, load balancing, etc.
 - Key escrow (e.g., reverse proxies for ssh, SSL)

Firewalls Pro/Con

- Benefits

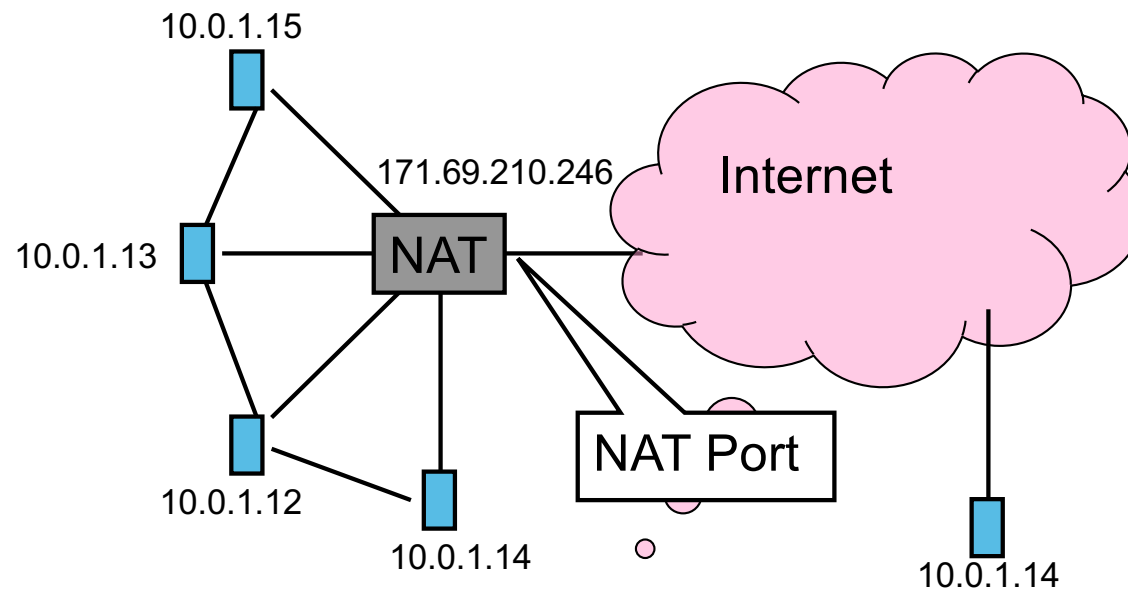
- Reduced “attack surface” against external attackers
- Filter out lots of “noise” in network traffic (helps focus attention)
- Reduced liability (common practice)

- Costs

- Actual cost: both hardware and administration
- Bottleneck and single point of failure on network
- False sense of security
 - Limited language (addresses, ports); doesn’t help with worms/viruses, ssh exploits, cross-site scripting, etc
 - Inside vs outside model is fragile (once an internal host is compromised firewall does no good); What about wireless laptops?
 - Modern companies increasingly offer no additional trust to machines inside the firewall (so-called “zero trust” architectures)

Network Address Translation

- Idea: Break the invariant that IP addresses are globally unique
 - Special addresses that are only **local**: 10.x.x.x, 192.168.x.x and 172.16.0.0-172.31.255.255



Typical NAT Behavior

- NAT maintains a table of the form:
 <client IP> <client port> <NAT ID>
- Outgoing packets (on internal port):
 - Look for client IP address, client port in the mapping table
 - If found, replace client port with previously allocated NAT ID (same size as port #)
 - If not found, allocate a new unique NAT ID and replace source port with NAT ID
 - Replace source address with NAT address (i.e., public IP address)
- Incoming Packets (on NAT port)
 - Look up destination port number as NAT ID in port mapping table
 - If found, replace destination address and port with client entries from the mapping table
 - If not found, the packet is not for us and should be rejected
- Table entries expire after 2-3 minutes of no activity to allow them to be garbage collected

NAT Pro/Con

- Benefits

- Only allows connections to the outside that are established from *inside*.
 - Hosts from outside can only contact internal hosts that appear in the mapping table, and they're only added when they establish the connection
- Don't need as large an external address space
 - (e.g., 10 machines can share 1 IP address)

- Costs

- Rewriting IP addresses isn't always easy (what if they appear in the **content** of the packet too? e.g., FTP. Then what happens to sequence numbers?)
- Breaks some protocols (e.g., some streaming protocols have client invoke server and then server opens new connection to client)
- But we've paid these costs, NAT is now ubiquitous...

Network content analysis

- Lots of devices want to look at network traffic **content** for security
 - Network Intrusion detection/prevention Systems (NIDS/NIPS)
 - Try to find signatures of attacks or malware
 - Spam filters
 - Try to detect unwanted e-mail
 - Data leakage
 - Try to prevent sensitive information from leaving company
 - Traffic differentiation
 - Filter or slow down BitTorrent traffic, Netflix traffic, etc
- Doing this as in the network is attractive because its cheaper and easier to manage than putting endpoint monitoring on each host

Challenges


- Expensive to look into each packet
 - 10Gbps -> ~1M packets per second... ns' per byte
 - Also must reassemble pkts into in-order streams (e.g., what if signature you are looking for spans two packets)
- Network vantage point is imperfect
 - What does a packet *mean*?
 - What if a session is encrypted?
 - Network evasion?

Network evasion


- Typically network intrusion detection systems are deployed like firewalls (between internal network and Internet)
- Key assumption is that NIDS sees the same traffic as destination host
- Not quite true...
 - Lots of ways to evade a NIDS by exploiting ambiguity

TTL evasion

- Suppose destination host is **2 hops** inside network *after* NIDS box
- This is what NIDS sees (each letter is one pkt)




- This is what the destination host sees



Sequence # evasion

- Suppose attacker sends two packets



- What does destination see?
 - IAMBED? IAMBAD? IAMBBD? Depends on host
- **Lots** of other evasion techniques...

Solution

- Protocol normalization
 - NIDS rewrites packets to remove all ambiguity
 - E.g.
 - all packets have ttl rewritten to reach any internal destination
 - IDS tracks each flow and does not allow overlapping packets
- Can be very tricky to get right and expensive
 - Potentially must buffer large amounts of data
 - What if you get seq #2 through #100, but not seq #1?
 - Tradeoff: when to drop data vs when to buffer

Bottom line

- The network vantage point is a very appealing place to implement policy because its central – **everyone has to go through it!**
 - But very challenging to infer the semantics of unknown communication
 - Port numbers communication very limited info
 - Parsing content is tough and many ways to evade
 - End-to-end encryption makes it hard to get content
- In spite of this, everyone uses these approaches because they don't have anything better

Next time

- User authentication